

# Build a Regression Model in Keras \_ Part A

## Result

`predictors.mean()`

Item	Mean
Cement	281.167864
Blast Furnace Slag	73.895825
Fly Ash	54.188350
Water	181.567282
Superplasticizer	6.204660
Coarse Aggregate	972.918932
Fine Aggregate	773.580485
Age	45.662136

`dtype: float64`

`predictors.std()`

Item	Mean
Cement	104.506364
Blast Furnace Slag	86.279342
Fly Ash	63.997004
Water	21.354219
Superplasticizer	5.973841
Coarse Aggregate	77.753954
Fine Aggregate	80.175980
Age	63.169912

`dtype: float64`

## Table of Contents

1. [Download and Clean Dataset](#) 2. [Import Keras](#) 3. [Build a Neural Network](#) 4. [Train and Test the Network](#)

# Download and Clean Dataset

Let's start by importing the *pandas* and the Numpy libraries.

In [1]:

```
import pandas as pd
import numpy as np
```

We will be playing around with the same dataset that we used in the videos.

The dataset is about the compressive strength of different samples of concrete based on the volumes of the different ingredients that were used to make them. Ingredients include:

1. Cement
2. Blast Furnace Slag
3. Fly Ash
4. Water
5. Superplasticizer
6. Coarse Aggregate
7. Fine Aggregate

Let's download the data and read it into a *pandas* dataframe.

In [2]:

```
concrete_data = pd.read_csv('https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0101EN/labs/data/concrete_data.csv')
concrete_data.head()
```

Out[2]:

	Cement	Blast Furnace Slag	Fly Ash	Water	Superplasticizer	Coarse Aggregate	Fine Aggregate	Age	Strength
0	540.0	0.0	0.0	162.0	2.5	1040.0	676.0	28	79.99
1	540.0	0.0	0.0	162.0	2.5	1055.0	676.0	28	61.89
2	332.5	142.5	0.0	228.0	0.0	932.0	594.0	270	40.27
3	332.5	142.5	0.0	228.0	0.0	932.0	594.0	365	41.05
4	198.6	132.4	0.0	192.0	0.0	978.4	825.5	360	44.30

So the first concrete sample has 540 cubic meter of cement, 0 cubic meter of blast furnace slag, 0 cubic meter of fly ash, 162 cubic meter of water, 2.5 cubic meter of superplasticizer, 1040 cubic meter of coarse aggregate, 676 cubic meter of fine aggregate. Such a concrete mix which is 28 days old, has a compressive strength of 79.99 MPa.

## Split data into predictors and target

The target variable in this problem is the concrete sample strength. Therefore, our predictors will be all the other columns.

In [3]:

```
concrete_data_columns = concrete_data.columns

predictors = concrete_data[concrete_data_columns[concrete_data_columns != 'Strength']] # all columns except Strength
target = concrete_data['Strength'] # Strength column
```

Let's do a quick sanity check of the predictors and the target dataframes.

Finally, the last step is to normalize the data by subtracting the mean and dividing by the standard deviation.

Let's save the number of predictors to *n\_cols* since we will need this number when building our network.

In [4]:

```
n_cols = predictors.shape[1] # number of predictors
```

## Import Keras

Recall from the videos that Keras normally runs on top of a low-level library such as TensorFlow. This means that to be able to use the Keras library, you will have to install TensorFlow first and when you import the Keras library, it will be explicitly displayed what backend was used to install the Keras library. In CC Labs, we used TensorFlow as the backend to install Keras, so it should clearly print that when we import Keras.

**Let's go ahead and import the Keras library**

In [5]:

```
import keras
```

Using TensorFlow backend.

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:521: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:522: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:523: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint32 = np.dtype [("qint32", np.int32, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:528: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
np_resource = np.dtype [("resource", np.ubyte, 1)]
```

As you can see, the TensorFlow backend was used to install the Keras library.

Let's import the rest of the packages from the Keras library that we will need to build our regression model.

In [6]:

```
from keras.models import Sequential
from keras.layers import Dense
```

## Build a Neural Network

Let's define a function that defines our regression model for us so that we can conveniently call it to create our model.

In [7]:

```
# define regression model
def regression_model():
    # create model
    model = Sequential()
    model.add(Dense(10, activation='relu', input_shape=(n_cols,)))
    model.add(Dense(1))

    # compile model
    model.compile(optimizer='adam', loss='mean_squared_error')
    return model
```

The above function create a model that has two hidden layers, each of 10 hidden units.

## Train and Test the Network

Let's call the function now to create our model.

In [8]:

```
# build the model
model = regression_model()
```

Next, we will train and test the model at the same time using the *fit* method. We will leave out 30% of the data for validation and we will train the model for 50 epochs.

In [9]:

```
# fit the model  
model.fit(predictors, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50  
- 1s - loss: 40562.6075 - val\_loss: 19354.4713

Epoch 2/50  
- 0s - loss: 8997.8915 - val\_loss: 5446.0898

Epoch 3/50  
- 0s - loss: 3408.7361 - val\_loss: 3742.4317

Epoch 4/50  
- 0s - loss: 2916.0421 - val\_loss: 3533.7680

Epoch 5/50  
- 0s - loss: 2695.0247 - val\_loss: 3413.3553

Epoch 6/50  
- 0s - loss: 2482.8500 - val\_loss: 3269.1465

Epoch 7/50  
- 0s - loss: 2289.8921 - val\_loss: 3106.4426

Epoch 8/50  
- 0s - loss: 2102.1340 - val\_loss: 2951.4280

Epoch 9/50  
- 0s - loss: 1925.2430 - val\_loss: 2804.2327

Epoch 10/50  
- 0s - loss: 1769.0078 - val\_loss: 2677.3645

Epoch 11/50  
- 0s - loss: 1627.3966 - val\_loss: 2559.1774

Epoch 12/50  
- 0s - loss: 1504.8699 - val\_loss: 2443.4871

Epoch 13/50  
- 0s - loss: 1383.9740 - val\_loss: 2312.0766

Epoch 14/50  
- 0s - loss: 1285.6321 - val\_loss: 2224.8974

Epoch 15/50  
- 0s - loss: 1197.4286 - val\_loss: 2145.6935

Epoch 16/50  
- 0s - loss: 1117.1521 - val\_loss: 2033.2754

Epoch 17/50  
- 0s - loss: 1048.5382 - val\_loss: 1965.0139

Epoch 18/50  
- 0s - loss: 988.5790 - val\_loss: 1878.0319

Epoch 19/50  
- 0s - loss: 929.8005 - val\_loss: 1807.1395

Epoch 20/50  
- 0s - loss: 882.2844 - val\_loss: 1732.6019

Epoch 21/50  
- 0s - loss: 839.3709 - val\_loss: 1662.2680

Epoch 22/50  
- 0s - loss: 799.5471 - val\_loss: 1604.6183

Epoch 23/50  
- 0s - loss: 765.7856 - val\_loss: 1550.7620

Epoch 24/50  
- 0s - loss: 729.9290 - val\_loss: 1491.4123

Epoch 25/50  
- 0s - loss: 702.2763 - val\_loss: 1437.7337

Epoch 26/50  
- 0s - loss: 675.4987 - val\_loss: 1381.2722

Epoch 27/50  
- 0s - loss: 649.3208 - val\_loss: 1336.0414

Epoch 28/50  
- 0s - loss: 627.0339 - val\_loss: 1284.3306

```
Epoch 29/50
- 0s - loss: 605.3076 - val_loss: 1240.0054
Epoch 30/50
- 0s - loss: 584.6712 - val_loss: 1193.6695
Epoch 31/50
- 0s - loss: 565.0752 - val_loss: 1155.4967
Epoch 32/50
- 0s - loss: 547.0044 - val_loss: 1108.0508
Epoch 33/50
- 0s - loss: 536.5699 - val_loss: 1071.7931
Epoch 34/50
- 0s - loss: 511.5580 - val_loss: 1024.1752
Epoch 35/50
- 0s - loss: 495.3470 - val_loss: 978.8584
Epoch 36/50
- 0s - loss: 479.5066 - val_loss: 948.9949
Epoch 37/50
- 0s - loss: 465.4431 - val_loss: 900.5298
Epoch 38/50
- 0s - loss: 450.0316 - val_loss: 875.6420
Epoch 39/50
- 0s - loss: 436.2808 - val_loss: 839.7511
Epoch 40/50
- 0s - loss: 423.4676 - val_loss: 800.7503
Epoch 41/50
- 0s - loss: 411.6116 - val_loss: 780.8262
Epoch 42/50
- 0s - loss: 398.5978 - val_loss: 731.8497
Epoch 43/50
- 0s - loss: 386.3814 - val_loss: 714.4407
Epoch 44/50
- 0s - loss: 375.2968 - val_loss: 660.4929
Epoch 45/50
- 0s - loss: 361.4536 - val_loss: 646.0115
Epoch 46/50
- 0s - loss: 353.6602 - val_loss: 605.1034
Epoch 47/50
- 0s - loss: 342.0523 - val_loss: 578.3189
Epoch 48/50
- 0s - loss: 334.0425 - val_loss: 554.6108
Epoch 49/50
- 0s - loss: 322.5900 - val_loss: 534.8040
Epoch 50/50
- 0s - loss: 313.2981 - val_loss: 509.0243
```

Out[9]:

```
<keras.callbacks.History at 0x7faa1ef06588>
```

**You can refer to this [link](<https://keras.io/models/sequential/>) to learn about other functions that you can use for prediction or evaluation.**

In [10]:

```
model = regression_model()
```



In [11]:

```
model.fit(predictors, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

- 1s - loss: 8416.2898 - val\_loss: 2930.2769

Epoch 2/50

- 0s - loss: 3852.0342 - val\_loss: 2420.4482

Epoch 3/50

- 0s - loss: 3313.5255 - val\_loss: 2155.0027

Epoch 4/50

- 0s - loss: 2942.2414 - val\_loss: 1841.9167

Epoch 5/50

- 0s - loss: 2596.5308 - val\_loss: 1629.2634

Epoch 6/50

- 0s - loss: 2280.7058 - val\_loss: 1460.3179

Epoch 7/50

- 0s - loss: 1995.0305 - val\_loss: 1299.6314

Epoch 8/50

- 0s - loss: 1748.1031 - val\_loss: 1162.1754

Epoch 9/50

- 0s - loss: 1528.8826 - val\_loss: 1055.3529

Epoch 10/50

- 0s - loss: 1344.5965 - val\_loss: 981.1630

Epoch 11/50

- 0s - loss: 1172.9462 - val\_loss: 894.0042

Epoch 12/50

- 0s - loss: 1031.0918 - val\_loss: 838.3811

Epoch 13/50

- 0s - loss: 911.7652 - val\_loss: 786.0792

Epoch 14/50

- 0s - loss: 812.3485 - val\_loss: 762.5485

Epoch 15/50

- 0s - loss: 722.2308 - val\_loss: 714.5188

Epoch 16/50

- 0s - loss: 645.6830 - val\_loss: 694.8396

Epoch 17/50

- 0s - loss: 584.3922 - val\_loss: 664.5404

Epoch 18/50

- 0s - loss: 530.5072 - val\_loss: 649.6963

Epoch 19/50

- 0s - loss: 492.0117 - val\_loss: 618.6195

Epoch 20/50

- 0s - loss: 446.1357 - val\_loss: 598.7580

Epoch 21/50

- 0s - loss: 411.2005 - val\_loss: 586.2897

Epoch 22/50

- 0s - loss: 385.8526 - val\_loss: 557.4580

Epoch 23/50

- 0s - loss: 358.8177 - val\_loss: 536.8256

Epoch 24/50

- 0s - loss: 338.1012 - val\_loss: 514.2364

Epoch 25/50

- 0s - loss: 318.0001 - val\_loss: 496.0567

Epoch 26/50

- 0s - loss: 301.4869 - val\_loss: 470.1007

Epoch 27/50

- 0s - loss: 288.5329 - val\_loss: 444.3394

Epoch 28/50

- 0s - loss: 272.6623 - val\_loss: 422.8754

```
Epoch 29/50
- 0s - loss: 260.8902 - val_loss: 402.3423
Epoch 30/50
- 0s - loss: 246.8919 - val_loss: 379.0095
Epoch 31/50
- 0s - loss: 236.4207 - val_loss: 358.1976
Epoch 32/50
- 0s - loss: 226.6401 - val_loss: 351.0278
Epoch 33/50
- 0s - loss: 219.0350 - val_loss: 318.1313
Epoch 34/50
- 0s - loss: 209.6482 - val_loss: 306.6366
Epoch 35/50
- 0s - loss: 201.5324 - val_loss: 283.7971
Epoch 36/50
- 0s - loss: 193.8761 - val_loss: 273.6053
Epoch 37/50
- 0s - loss: 187.4440 - val_loss: 251.6677
Epoch 38/50
- 0s - loss: 181.8609 - val_loss: 238.5323
Epoch 39/50
- 0s - loss: 176.9942 - val_loss: 225.8732
Epoch 40/50
- 0s - loss: 171.5999 - val_loss: 216.0015
Epoch 41/50
- 0s - loss: 166.1838 - val_loss: 220.8213
Epoch 42/50
- 0s - loss: 163.1542 - val_loss: 196.0382
Epoch 43/50
- 0s - loss: 159.5071 - val_loss: 187.7536
Epoch 44/50
- 0s - loss: 155.7434 - val_loss: 179.4457
Epoch 45/50
- 0s - loss: 152.6399 - val_loss: 171.4061
Epoch 46/50
- 0s - loss: 150.1296 - val_loss: 166.2962
Epoch 47/50
- 0s - loss: 147.0033 - val_loss: 161.4621
Epoch 48/50
- 0s - loss: 146.7481 - val_loss: 155.1636
Epoch 49/50
- 0s - loss: 142.6591 - val_loss: 149.1615
Epoch 50/50
- 0s - loss: 143.0025 - val_loss: 145.4346
```

Out[11]:

<keras.callbacks.History at 0x7faalc1e28d0>

In [12]:

```
model = regression_model()
```

In [13]:

```
model.fit(predictors, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

- 1s - loss: 795.2996 - val\_loss: 554.0186

Epoch 2/50

- 0s - loss: 523.5439 - val\_loss: 470.9915

Epoch 3/50

- 0s - loss: 416.7169 - val\_loss: 307.7284

Epoch 4/50

- 0s - loss: 346.4876 - val\_loss: 276.3472

Epoch 5/50

- 0s - loss: 318.4829 - val\_loss: 226.1263

Epoch 6/50

- 0s - loss: 307.7949 - val\_loss: 213.1514

Epoch 7/50

- 0s - loss: 298.6301 - val\_loss: 191.7274

Epoch 8/50

- 0s - loss: 292.3293 - val\_loss: 188.5618

Epoch 9/50

- 0s - loss: 287.0091 - val\_loss: 175.1596

Epoch 10/50

- 0s - loss: 281.6690 - val\_loss: 165.7769

Epoch 11/50

- 0s - loss: 277.3691 - val\_loss: 161.5906

Epoch 12/50

- 0s - loss: 272.4128 - val\_loss: 159.5739

Epoch 13/50

- 0s - loss: 268.0451 - val\_loss: 154.3427

Epoch 14/50

- 0s - loss: 264.3309 - val\_loss: 150.1205

Epoch 15/50

- 0s - loss: 262.4972 - val\_loss: 147.3556

Epoch 16/50

- 0s - loss: 258.0879 - val\_loss: 145.7099

Epoch 17/50

- 0s - loss: 254.7169 - val\_loss: 145.9030

Epoch 18/50

- 0s - loss: 251.9108 - val\_loss: 142.5226

Epoch 19/50

- 0s - loss: 249.2911 - val\_loss: 142.6395

Epoch 20/50

- 0s - loss: 246.3667 - val\_loss: 140.1178

Epoch 21/50

- 0s - loss: 243.4286 - val\_loss: 140.2505

Epoch 22/50

- 0s - loss: 240.7583 - val\_loss: 138.3024

Epoch 23/50

- 0s - loss: 238.0386 - val\_loss: 137.6034

Epoch 24/50

- 0s - loss: 235.8715 - val\_loss: 133.9560

Epoch 25/50

- 0s - loss: 232.5535 - val\_loss: 133.9363

Epoch 26/50

- 0s - loss: 229.8911 - val\_loss: 131.4385

Epoch 27/50

- 0s - loss: 227.2344 - val\_loss: 129.2517

Epoch 28/50

- 0s - loss: 224.6342 - val\_loss: 128.4986

```
Epoch 29/50
- 0s - loss: 221.5578 - val_loss: 126.2728
Epoch 30/50
- 0s - loss: 219.0286 - val_loss: 124.1832
Epoch 31/50
- 0s - loss: 217.2474 - val_loss: 124.0830
Epoch 32/50
- 0s - loss: 214.0404 - val_loss: 121.9254
Epoch 33/50
- 0s - loss: 210.8602 - val_loss: 121.2150
Epoch 34/50
- 0s - loss: 208.6816 - val_loss: 118.5277
Epoch 35/50
- 0s - loss: 205.4320 - val_loss: 117.8958
Epoch 36/50
- 0s - loss: 203.4639 - val_loss: 113.3904
Epoch 37/50
- 0s - loss: 200.3719 - val_loss: 117.9152
Epoch 38/50
- 0s - loss: 197.6920 - val_loss: 109.3236
Epoch 39/50
- 0s - loss: 195.3348 - val_loss: 111.8046
Epoch 40/50
- 0s - loss: 193.0749 - val_loss: 108.9683
Epoch 41/50
- 0s - loss: 190.3472 - val_loss: 105.7171
Epoch 42/50
- 0s - loss: 188.1152 - val_loss: 105.9286
Epoch 43/50
- 1s - loss: 185.5927 - val_loss: 104.4759
Epoch 44/50
- 0s - loss: 183.3283 - val_loss: 103.1977
Epoch 45/50
- 0s - loss: 181.3317 - val_loss: 101.8796
Epoch 46/50
- 0s - loss: 178.7418 - val_loss: 100.0174
Epoch 47/50
- 0s - loss: 177.4148 - val_loss: 103.7187
Epoch 48/50
- 0s - loss: 174.8999 - val_loss: 97.1213
Epoch 49/50
- 0s - loss: 172.8278 - val_loss: 96.2261
Epoch 50/50
- 0s - loss: 170.6372 - val_loss: 98.3315
```

Out[13]:

```
<keras.callbacks.History at 0x7faa081c5240>
```

In [14]:

```
score = model.evaluate(predictors, target)
```

```
1030/1030 [=====] - 0s 99us/step
```

Feel free to vary the following and note what impact each change has on the model's performance:

1. Increase or decrease number of neurons in hidden layers
2. Add more hidden layers
3. Increase number of epochs

In [15]:

```
target.mean()
```

Out[15]:

```
35.817961165048544
```

In [16]:

```
target.std()
```

Out[16]:

```
16.705741961912512
```

In [17]:

```
predictors.mean()
```

Out[17]:

Cement	281.167864
Blast Furnace Slag	73.895825
Fly Ash	54.188350
Water	181.567282
Superplasticizer	6.204660
Coarse Aggregate	972.918932
Fine Aggregate	773.580485
Age	45.662136

dtype: float64

In [18]:

```
predictors.std()
```

Out[18]:

Cement	104.506364
Blast Furnace Slag	86.279342
Fly Ash	63.997004
Water	21.354219
Superplasticizer	5.973841
Coarse Aggregate	77.753954
Fine Aggregate	80.175980
Age	63.169912

dtype: float64

In [19]:

```
model.mean()
```

```
-----  
-----  
AttributeError                                Traceback (most recent call 1  
ast)  
<ipython-input-19-fd5124642809> in <module>  
----> 1 model.mean()  
  
AttributeError: 'Sequential' object has no attribute 'mean'
```

In [ ]: