

Build a Regression Model in Keras _ Part B

Result

predictors.mean()

Item	Mean
Cement	2.432224e-15
Blast Furnace Slag	-8.513686e-16
Fly Ash	3.837815e-16
Water	1.846743e-15
Superplasticizer	-9.641155e-16
Coarse Aggregate	6.818710e-15
Fine Aggregate	1.232571e-14
Age	3.640022e-16

dtype: float64

Table of Contents

1. [Download and Clean Dataset](#) 2. [Import Keras](#) 3. [Build a Neural Network](#) 4. [Train and Test the Network](#)

Download and Clean Dataset

Let's start by importing the *pandas* and the Numpy libraries.

In [2]:

```
import pandas as pd
import numpy as np
```

We will be playing around with the same dataset that we used in the videos.

The dataset is about the compressive strength of different samples of concrete based on the volumes of the different ingredients that were used to make them. Ingredients include:

1. Cement
2. Blast Furnace Slag
3. Fly Ash
4. Water
5. Superplasticizer
6. Coarse Aggregate
7. Fine Aggregate

Let's download the data and read it into a *pandas* dataframe.

In [3]:

```
concrete_data = pd.read_csv('https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0101EN/labs/data/concrete_data.csv')
concrete_data.head()
```

Out[3]:

	Cement	Blast Furnace Slag	Fly Ash	Water	Superplasticizer	Coarse Aggregate	Fine Aggregate	Age	Strength
0	540.0	0.0	0.0	162.0	2.5	1040.0	676.0	28	79.99
1	540.0	0.0	0.0	162.0	2.5	1055.0	676.0	28	61.89
2	332.5	142.5	0.0	228.0	0.0	932.0	594.0	270	40.27
3	332.5	142.5	0.0	228.0	0.0	932.0	594.0	365	41.05
4	198.6	132.4	0.0	192.0	0.0	978.4	825.5	360	44.30

So the first concrete sample has 540 cubic meter of cement, 0 cubic meter of blast furnace slag, 0 cubic meter of fly ash, 162 cubic meter of water, 2.5 cubic meter of superplasticizer, 1040 cubic meter of coarse aggregate, 676 cubic meter of fine aggregate. Such a concrete mix which is 28 days old, has a compressive strength of 79.99 MPa.

Split data into predictors and target

The target variable in this problem is the concrete sample strength. Therefore, our predictors will be all the other columns.

In [4]:

```
concrete_data_columns = concrete_data.columns

predictors = concrete_data[concrete_data_columns[concrete_data_columns != 'Strength']] # all columns except Strength
target = concrete_data['Strength'] # Strength column
```

Let's do a quick sanity check of the predictors and the target dataframes.

Finally, the last step is to normalize the data by subtracting the mean and dividing by the standard deviation.

In [5]:

```
predictors_norm = (predictors - predictors.mean()) / predictors.std()
predictors_norm.head()
```

Out[5]:

	Cement	Blast Furnace Slag	Fly Ash	Water	Superplasticizer	Coarse Aggregate	Fine Aggregate	Age
0	2.476712	-0.856472	-0.846733	-0.916319	-0.620147	0.862735	-1.217079	-0.279597
1	2.476712	-0.856472	-0.846733	-0.916319	-0.620147	1.055651	-1.217079	-0.279597
2	0.491187	0.795140	-0.846733	2.174405	-1.038638	-0.526262	-2.239829	3.551340
3	0.491187	0.795140	-0.846733	2.174405	-1.038638	-0.526262	-2.239829	5.055221
4	-0.790075	0.678079	-0.846733	0.488555	-1.038638	0.070492	0.647569	4.976069

Let's save the number of predictors to n_cols since we will need this number when building our network.

In [6]:

```
n_cols = predictors_norm.shape[1] # number of predictors
```

Import Keras

Recall from the videos that Keras normally runs on top of a low-level library such as TensorFlow. This means that to be able to use the Keras library, you will have to install TensorFlow first and when you import the Keras library, it will be explicitly displayed what backend was used to install the Keras library. In CC Labs, we used TensorFlow as the backend to install Keras, so it should clearly print that when we import Keras.

Let's go ahead and import the Keras library

In [7]:

```
import keras
```

Using TensorFlow backend.

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:521: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:522: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:523: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint32 = np.dtype [("qint32", np.int32, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:528: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
np_resource = np.dtype [("resource", np.ubyte, 1)]
```

As you can see, the TensorFlow backend was used to install the Keras library.

Let's import the rest of the packages from the Keras library that we will need to build our regression model.

In [8]:

```
from keras.models import Sequential
from keras.layers import Dense
```

Build a Neural Network

Let's define a function that defines our regression model for us so that we can conveniently call it to create our model.

In [9]:

```
# define regression model
def regression_model():
    # create model
    model = Sequential()
    model.add(Dense(10, activation='relu', input_shape=(n_cols,)))
    model.add(Dense(1))

    # compile model
    model.compile(optimizer='adam', loss='mean_squared_error')
    return model
```

The above function create a model that has two hidden layers, each of 10 hidden units.

Train and Test the Network

Let's call the function now to create our model.

In [10]:

```
# build the model
model = regression_model()
```

Next, we will train and test the model at the same time using the *fit* method. We will leave out 30% of the data for validation and we will train the model for 50 epochs.

In [11]:

```
# fit the model  
model.fit(predictors_norm, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

- 2s - loss: 1711.4989 - val_loss: 1221.8312

Epoch 2/50

- 0s - loss: 1692.4718 - val_loss: 1209.9213

Epoch 3/50

- 0s - loss: 1673.0860 - val_loss: 1197.9194

Epoch 4/50

- 1s - loss: 1653.3507 - val_loss: 1185.7700

Epoch 5/50

- 1s - loss: 1633.0977 - val_loss: 1173.5639

Epoch 6/50

- 1s - loss: 1611.7509 - val_loss: 1160.8569

Epoch 7/50

- 1s - loss: 1589.7819 - val_loss: 1147.9310

Epoch 8/50

- 1s - loss: 1566.4176 - val_loss: 1134.5165

Epoch 9/50

- 1s - loss: 1541.9591 - val_loss: 1120.4375

Epoch 10/50

- 1s - loss: 1516.4314 - val_loss: 1105.9220

Epoch 11/50

- 1s - loss: 1489.3917 - val_loss: 1091.0506

Epoch 12/50

- 0s - loss: 1461.6187 - val_loss: 1075.4174

Epoch 13/50

- 1s - loss: 1432.4808 - val_loss: 1059.5400

Epoch 14/50

- 0s - loss: 1402.4288 - val_loss: 1042.8070

Epoch 15/50

- 1s - loss: 1370.8199 - val_loss: 1025.9860

Epoch 16/50

- 1s - loss: 1339.1600 - val_loss: 1008.4214

Epoch 17/50

- 1s - loss: 1305.4521 - val_loss: 990.5801

Epoch 18/50

- 1s - loss: 1271.0235 - val_loss: 972.4215

Epoch 19/50

- 1s - loss: 1235.8714 - val_loss: 953.8644

Epoch 20/50

- 1s - loss: 1200.4883 - val_loss: 934.6821

Epoch 21/50

- 0s - loss: 1164.6588 - val_loss: 915.2722

Epoch 22/50

- 1s - loss: 1128.4606 - val_loss: 895.5838

Epoch 23/50

- 0s - loss: 1091.9728 - val_loss: 875.9351

Epoch 24/50

- 1s - loss: 1055.4538 - val_loss: 855.7119

Epoch 25/50

- 1s - loss: 1019.5089 - val_loss: 835.4262

Epoch 26/50

- 1s - loss: 983.3954 - val_loss: 814.9198

Epoch 27/50

- 0s - loss: 947.7911 - val_loss: 794.9194

Epoch 28/50

- 0s - loss: 912.3611 - val_loss: 774.8496

```
Epoch 29/50
- 0s - loss: 877.6398 - val_loss: 755.0444
Epoch 30/50
- 0s - loss: 843.9990 - val_loss: 734.8582
Epoch 31/50
- 0s - loss: 811.3524 - val_loss: 715.1335
Epoch 32/50
- 0s - loss: 779.5729 - val_loss: 695.0538
Epoch 33/50
- 0s - loss: 748.3396 - val_loss: 675.8676
Epoch 34/50
- 0s - loss: 719.0644 - val_loss: 656.0712
Epoch 35/50
- 0s - loss: 689.8357 - val_loss: 637.6149
Epoch 36/50
- 0s - loss: 662.3096 - val_loss: 619.1739
Epoch 37/50
- 0s - loss: 636.0754 - val_loss: 600.4880
Epoch 38/50
- 0s - loss: 610.2736 - val_loss: 582.4201
Epoch 39/50
- 0s - loss: 586.4063 - val_loss: 564.4967
Epoch 40/50
- 0s - loss: 563.2232 - val_loss: 547.2692
Epoch 41/50
- 0s - loss: 541.2252 - val_loss: 530.7381
Epoch 42/50
- 0s - loss: 520.6786 - val_loss: 514.3229
Epoch 43/50
- 0s - loss: 500.7559 - val_loss: 497.9842
Epoch 44/50
- 0s - loss: 482.0781 - val_loss: 482.3282
Epoch 45/50
- 0s - loss: 464.5343 - val_loss: 466.4495
Epoch 46/50
- 0s - loss: 447.7795 - val_loss: 451.4834
Epoch 47/50
- 0s - loss: 431.8689 - val_loss: 437.0190
Epoch 48/50
- 0s - loss: 417.0691 - val_loss: 423.2871
Epoch 49/50
- 0s - loss: 403.0833 - val_loss: 409.1368
Epoch 50/50
- 0s - loss: 389.5511 - val_loss: 395.7581
```

Out[11]:

```
<keras.callbacks.History at 0x7f0166367588>
```

You can refer to this [link](<https://keras.io/models/sequential/>) to learn about other functions that you can use for prediction or evaluation.

In [12]:

```
model = regression_model()
```


In [13]:

```
model.fit(predictors_norm, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

- 1s - loss: 1696.1677 - val_loss: 1205.1732

Epoch 2/50

- 0s - loss: 1674.4248 - val_loss: 1192.9089

Epoch 3/50

- 0s - loss: 1652.1294 - val_loss: 1180.3829

Epoch 4/50

- 0s - loss: 1629.5993 - val_loss: 1167.5001

Epoch 5/50

- 0s - loss: 1606.4634 - val_loss: 1154.0415

Epoch 6/50

- 0s - loss: 1582.5934 - val_loss: 1139.9782

Epoch 7/50

- 1s - loss: 1557.4204 - val_loss: 1125.3566

Epoch 8/50

- 0s - loss: 1531.4932 - val_loss: 1110.1046

Epoch 9/50

- 0s - loss: 1504.3119 - val_loss: 1094.3082

Epoch 10/50

- 0s - loss: 1475.7358 - val_loss: 1078.1465

Epoch 11/50

- 0s - loss: 1446.4053 - val_loss: 1061.3482

Epoch 12/50

- 1s - loss: 1415.9390 - val_loss: 1043.5264

Epoch 13/50

- 1s - loss: 1383.8973 - val_loss: 1025.4712

Epoch 14/50

- 1s - loss: 1351.6654 - val_loss: 1006.6807

Epoch 15/50

- 1s - loss: 1318.1523 - val_loss: 987.8780

Epoch 16/50

- 0s - loss: 1284.1083 - val_loss: 968.4411

Epoch 17/50

- 0s - loss: 1249.1506 - val_loss: 948.5207

Epoch 18/50

- 1s - loss: 1213.5545 - val_loss: 928.2050

Epoch 19/50

- 1s - loss: 1177.6338 - val_loss: 908.0533

Epoch 20/50

- 1s - loss: 1140.9328 - val_loss: 887.1632

Epoch 21/50

- 1s - loss: 1104.5342 - val_loss: 866.1762

Epoch 22/50

- 1s - loss: 1067.3289 - val_loss: 844.7585

Epoch 23/50

- 1s - loss: 1030.3784 - val_loss: 823.5787

Epoch 24/50

- 1s - loss: 993.4599 - val_loss: 802.8675

Epoch 25/50

- 0s - loss: 957.4142 - val_loss: 781.5620

Epoch 26/50

- 1s - loss: 921.0617 - val_loss: 760.7687

Epoch 27/50

- 1s - loss: 885.9281 - val_loss: 739.4633

Epoch 28/50

- 0s - loss: 850.5479 - val_loss: 718.9233

```
Epoch 29/50
- 1s - loss: 816.8905 - val_loss: 698.0400
Epoch 30/50
- 1s - loss: 783.7427 - val_loss: 677.2268
Epoch 31/50
- 1s - loss: 751.1527 - val_loss: 657.4656
Epoch 32/50
- 0s - loss: 720.2465 - val_loss: 637.8150
Epoch 33/50
- 0s - loss: 690.1726 - val_loss: 618.5611
Epoch 34/50
- 1s - loss: 660.7655 - val_loss: 600.2462
Epoch 35/50
- 1s - loss: 633.2927 - val_loss: 581.3520
Epoch 36/50
- 1s - loss: 605.8629 - val_loss: 563.6073
Epoch 37/50
- 1s - loss: 580.4696 - val_loss: 545.6735
Epoch 38/50
- 0s - loss: 555.7655 - val_loss: 528.6938
Epoch 39/50
- 1s - loss: 532.5409 - val_loss: 512.2938
Epoch 40/50
- 0s - loss: 510.4634 - val_loss: 496.6151
Epoch 41/50
- 1s - loss: 489.7261 - val_loss: 481.0181
Epoch 42/50
- 0s - loss: 470.1424 - val_loss: 466.7966
Epoch 43/50
- 0s - loss: 451.8457 - val_loss: 452.2396
Epoch 44/50
- 0s - loss: 434.3136 - val_loss: 439.2589
Epoch 45/50
- 1s - loss: 418.2094 - val_loss: 426.0502
Epoch 46/50
- 1s - loss: 403.1780 - val_loss: 413.8536
Epoch 47/50
- 1s - loss: 388.9723 - val_loss: 401.7463
Epoch 48/50
- 1s - loss: 375.5718 - val_loss: 391.1045
Epoch 49/50
- 1s - loss: 363.5033 - val_loss: 380.5180
Epoch 50/50
- 1s - loss: 352.1709 - val_loss: 370.6975
```

Out[13]:

<keras.callbacks.History at 0x7f0164551ef0>

In [14]:

```
model = regression_model()
```

In [15]:

```
model.fit(predictors_norm, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

- 2s - loss: 1728.5768 - val_loss: 1247.3424

Epoch 2/50

- 1s - loss: 1708.5013 - val_loss: 1236.0506

Epoch 3/50

- 0s - loss: 1688.9116 - val_loss: 1224.9578

Epoch 4/50

- 0s - loss: 1669.2916 - val_loss: 1214.4407

Epoch 5/50

- 1s - loss: 1649.7405 - val_loss: 1204.0223

Epoch 6/50

- 0s - loss: 1629.9785 - val_loss: 1193.6450

Epoch 7/50

- 1s - loss: 1609.9223 - val_loss: 1183.4566

Epoch 8/50

- 1s - loss: 1589.6233 - val_loss: 1173.2516

Epoch 9/50

- 1s - loss: 1569.4277 - val_loss: 1162.8928

Epoch 10/50

- 1s - loss: 1548.1741 - val_loss: 1152.7625

Epoch 11/50

- 1s - loss: 1526.9603 - val_loss: 1142.1318

Epoch 12/50

- 1s - loss: 1504.6511 - val_loss: 1131.4254

Epoch 13/50

- 0s - loss: 1482.2491 - val_loss: 1120.3812

Epoch 14/50

- 0s - loss: 1458.9476 - val_loss: 1109.1980

Epoch 15/50

- 1s - loss: 1435.2923 - val_loss: 1097.6263

Epoch 16/50

- 0s - loss: 1410.9882 - val_loss: 1085.8319

Epoch 17/50

- 1s - loss: 1386.1910 - val_loss: 1074.1326

Epoch 18/50

- 0s - loss: 1361.3060 - val_loss: 1061.9616

Epoch 19/50

- 0s - loss: 1335.8924 - val_loss: 1049.6028

Epoch 20/50

- 0s - loss: 1310.0226 - val_loss: 1036.8008

Epoch 21/50

- 1s - loss: 1283.8813 - val_loss: 1024.1250

Epoch 22/50

- 1s - loss: 1257.4994 - val_loss: 1010.8475

Epoch 23/50

- 1s - loss: 1230.6285 - val_loss: 997.4467

Epoch 24/50

- 1s - loss: 1203.8660 - val_loss: 983.8288

Epoch 25/50

- 1s - loss: 1177.0770 - val_loss: 969.5986

Epoch 26/50

- 1s - loss: 1149.7590 - val_loss: 954.7330

Epoch 27/50

- 0s - loss: 1122.0375 - val_loss: 939.8554

Epoch 28/50

- 1s - loss: 1094.1217 - val_loss: 924.8475

```
Epoch 29/50
- 1s - loss: 1065.8663 - val_loss: 908.1475
Epoch 30/50
- 0s - loss: 1036.7879 - val_loss: 890.9676
Epoch 31/50
- 1s - loss: 1007.5199 - val_loss: 873.1516
Epoch 32/50
- 1s - loss: 977.6479 - val_loss: 854.8861
Epoch 33/50
- 1s - loss: 946.8845 - val_loss: 836.1368
Epoch 34/50
- 1s - loss: 916.3753 - val_loss: 815.6918
Epoch 35/50
- 6s - loss: 884.4618 - val_loss: 794.5632
Epoch 36/50
- 0s - loss: 852.3296 - val_loss: 773.0929
Epoch 37/50
- 1s - loss: 820.0540 - val_loss: 749.8852
Epoch 38/50
- 1s - loss: 787.0309 - val_loss: 726.6136
Epoch 39/50
- 1s - loss: 753.7467 - val_loss: 702.8593
Epoch 40/50
- 0s - loss: 721.0309 - val_loss: 678.6941
Epoch 41/50
- 0s - loss: 688.7194 - val_loss: 653.6791
Epoch 42/50
- 1s - loss: 656.5256 - val_loss: 628.8406
Epoch 43/50
- 0s - loss: 625.3068 - val_loss: 603.5848
Epoch 44/50
- 1s - loss: 594.5003 - val_loss: 579.5138
Epoch 45/50
- 1s - loss: 564.8973 - val_loss: 555.7843
Epoch 46/50
- 1s - loss: 536.8039 - val_loss: 532.1986
Epoch 47/50
- 1s - loss: 509.7000 - val_loss: 510.3199
Epoch 48/50
- 1s - loss: 484.4889 - val_loss: 488.9479
Epoch 49/50
- 1s - loss: 460.3913 - val_loss: 468.0664
Epoch 50/50
- 0s - loss: 437.8351 - val_loss: 448.4843
```

Out[15]:

<keras.callbacks.History at 0x7f015473ba58>

In [16]:

```
score = model.evaluate(predictors_norm, target)
```

1030/1030 [=====] - 0s 115us/step

Feel free to vary the following and note what impact each change has on the model's performance:

1. Increase or decrease number of neurons in hidden layers
2. Add more hidden layers
3. Increase number of epochs

In [17]:

```
target.mean()
```

Out[17]:

```
35.817961165048544
```

In [18]:

```
target.std()
```

Out[18]:

```
16.705741961912512
```

In [20]:

```
predictors_norm.mean()
```

Out[20]:

```
Cement          2.432224e-15
Blast Furnace Slag -8.513686e-16
Fly Ash         3.837815e-16
Water           1.846743e-15
Superplasticizer -9.641155e-16
Coarse Aggregate 6.818710e-15
Fine Aggregate   1.232571e-14
Age             3.640022e-16
dtype: float64
```

In [21]:

```
predictors_norm.std()
```

Out[21]:

```
Cement          1.0
Blast Furnace Slag 1.0
Fly Ash         1.0
Water           1.0
Superplasticizer 1.0
Coarse Aggregate 1.0
Fine Aggregate   1.0
Age             1.0
dtype: float64
```

In []: