

Build a Regression Model in Keras _ Part B

Result

predictors.mean()

Item	Mean
Cement	2.432224e-15
Blast Furnace Slag	-8.513686e-16
Fly Ash	3.837815e-16
Water	1.846743e-15
Superplasticizer	-9.641155e-16
Coarse Aggregate	6.818710e-15
Fine Aggregate	1.232571e-14
Age	3.640022e-16

dtype: float64

Table of Contents

1. [Download and Clean Dataset](#) 2. [Import Keras](#) 3. [Build a Neural Network](#) 4. [Train and Test the Network](#)

Download and Clean Dataset

Let's start by importing the *pandas* and the Numpy libraries.

In [1]:

```
import pandas as pd
import numpy as np
```

We will be playing around with the same dataset that we used in the videos.

The dataset is about the compressive strength of different samples of concrete based on the volumes of the different ingredients that were used to make them. Ingredients include:

- 1. Cement**
- 2. Blast Furnace Slag**
- 3. Fly Ash**
- 4. Water**
- 5. Superplasticizer**
- 6. Coarse Aggregate**
- 7. Fine Aggregate**

Let's download the data and read it into a *pandas* dataframe.

In [2]:

```
concrete_data = pd.read_csv('https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0101EN/labs/data/concrete_data.csv')
concrete_data.head()
```

Out[2]:

	Cement	Blast Furnace Slag	Fly Ash	Water	Superplasticizer	Coarse Aggregate	Fine Aggregate	Age	Strength
0	540.0	0.0	0.0	162.0	2.5	1040.0	676.0	28	79.99
1	540.0	0.0	0.0	162.0	2.5	1055.0	676.0	28	61.89
2	332.5	142.5	0.0	228.0	0.0	932.0	594.0	270	40.27
3	332.5	142.5	0.0	228.0	0.0	932.0	594.0	365	41.05
4	198.6	132.4	0.0	192.0	0.0	978.4	825.5	360	44.30

So the first concrete sample has 540 cubic meter of cement, 0 cubic meter of blast furnace slag, 0 cubic meter of fly ash, 162 cubic meter of water, 2.5 cubic meter of superplasticizer, 1040 cubic meter of coarse aggregate, 676 cubic meter of fine aggregate. Such a concrete mix which is 28 days old, has a compressive strength of 79.99 MPa.

Split data into predictors and target

The target variable in this problem is the concrete sample strength. Therefore, our predictors will be all the other columns.

In [3]:

```
concrete_data_columns = concrete_data.columns

predictors = concrete_data[concrete_data_columns[concrete_data_columns != 'Strength']] # all columns except Strength
target = concrete_data['Strength'] # Strength column
```

Let's do a quick sanity check of the predictors and the target dataframes.

Finally, the last step is to normalize the data by subtracting the mean and dividing by the standard deviation.

In [4]:

```
predictors_norm = (predictors - predictors.mean()) / predictors.std()
predictors_norm.head()
```

Out[4]:

	Cement	Blast Furnace Slag	Fly Ash	Water	Superplasticizer	Coarse Aggregate	Fine Aggregate	Age
0	2.476712	-0.856472	-0.846733	-0.916319	-0.620147	0.862735	-1.217079	-0.279597
1	2.476712	-0.856472	-0.846733	-0.916319	-0.620147	1.055651	-1.217079	-0.279597
2	0.491187	0.795140	-0.846733	2.174405	-1.038638	-0.526262	-2.239829	3.551340
3	0.491187	0.795140	-0.846733	2.174405	-1.038638	-0.526262	-2.239829	5.055221
4	-0.790075	0.678079	-0.846733	0.488555	-1.038638	0.070492	0.647569	4.976069

Let's save the number of predictors to n_cols since we will need this number when building our network.

In [5]:

```
n_cols = predictors_norm.shape[1] # number of predictors
```

Import Keras

Recall from the videos that Keras normally runs on top of a low-level library such as TensorFlow. This means that to be able to use the Keras library, you will have to install TensorFlow first and when you import the Keras library, it will be explicitly displayed what backend was used to install the Keras library. In CC Labs, we used TensorFlow as the backend to install Keras, so it should clearly print that when we import Keras.

Let's go ahead and import the Keras library

In [6]:

```
import keras
```

Using TensorFlow backend.

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:521: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:522: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:523: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint32 = np.dtype [("qint32", np.int32, 1)]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:528: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
np_resource = np.dtype [("resource", np.ubyte, 1)]
```

As you can see, the TensorFlow backend was used to install the Keras library.

Let's import the rest of the packages from the Keras library that we will need to build our regression model.

In [7]:

```
from keras.models import Sequential
from keras.layers import Dense
```

Build a Neural Network

Let's define a function that defines our regression model for us so that we can conveniently call it to create our model.

In [8]:

```
# define regression model
def regression_model():
    # create model
    model = Sequential()
    model.add(Dense(10, activation='relu', input_shape=(n_cols,)))
    model.add(Dense(10, activation='relu'))
    model.add(Dense(10, activation='relu'))
    model.add(Dense(1))

    # compile model
    model.compile(optimizer='adam', loss='mean_squared_error')
    return model
```

The above function create a model that has two hidden layers, each of 10 hidden units.

Train and Test the Network

Let's call the function now to create our model.

In [9]:

```
# build the model
model = regression_model()
```

Next, we will train and test the model at the same time using the *fit* method. We will leave out 30% of the data for validation and we will train the model for 50 epochs.

In [10]:

```
# fit the model  
model.fit(predictors_norm, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

- 1s - loss: 1697.7120 - val_loss: 1224.8229

Epoch 2/50

- 1s - loss: 1680.1567 - val_loss: 1212.7504

Epoch 3/50

- 0s - loss: 1654.9304 - val_loss: 1192.4209

Epoch 4/50

- 0s - loss: 1614.7245 - val_loss: 1158.8317

Epoch 5/50

- 0s - loss: 1553.2393 - val_loss: 1108.7484

Epoch 6/50

- 0s - loss: 1464.8055 - val_loss: 1036.9128

Epoch 7/50

- 0s - loss: 1340.2559 - val_loss: 933.5972

Epoch 8/50

- 0s - loss: 1164.7225 - val_loss: 791.4932

Epoch 9/50

- 0s - loss: 938.6757 - val_loss: 628.6110

Epoch 10/50

- 0s - loss: 693.5522 - val_loss: 478.8902

Epoch 11/50

- 0s - loss: 480.2239 - val_loss: 378.8241

Epoch 12/50

- 0s - loss: 345.8603 - val_loss: 321.3862

Epoch 13/50

- 0s - loss: 276.8219 - val_loss: 286.5745

Epoch 14/50

- 0s - loss: 244.0334 - val_loss: 264.5930

Epoch 15/50

- 0s - loss: 228.8953 - val_loss: 250.1560

Epoch 16/50

- 0s - loss: 219.5803 - val_loss: 237.6332

Epoch 17/50

- 0s - loss: 212.8663 - val_loss: 231.3798

Epoch 18/50

- 0s - loss: 206.6765 - val_loss: 225.3033

Epoch 19/50

- 0s - loss: 202.5399 - val_loss: 220.5405

Epoch 20/50

- 0s - loss: 197.9348 - val_loss: 214.3778

Epoch 21/50

- 0s - loss: 194.1849 - val_loss: 210.1460

Epoch 22/50

- 0s - loss: 190.7083 - val_loss: 204.2309

Epoch 23/50

- 0s - loss: 187.4776 - val_loss: 198.0577

Epoch 24/50

- 0s - loss: 184.2677 - val_loss: 194.4762

Epoch 25/50

- 0s - loss: 181.0198 - val_loss: 191.0467

Epoch 26/50

- 0s - loss: 178.1203 - val_loss: 184.9691

Epoch 27/50

- 0s - loss: 175.3959 - val_loss: 181.3125

Epoch 28/50

- 1s - loss: 172.0204 - val_loss: 177.6041

```
Epoch 29/50
- 0s - loss: 169.4722 - val_loss: 173.5333
Epoch 30/50
- 0s - loss: 166.6460 - val_loss: 170.8332
Epoch 31/50
- 0s - loss: 164.1814 - val_loss: 167.6179
Epoch 32/50
- 0s - loss: 161.4800 - val_loss: 162.4757
Epoch 33/50
- 0s - loss: 158.7780 - val_loss: 159.4576
Epoch 34/50
- 0s - loss: 155.9600 - val_loss: 154.1154
Epoch 35/50
- 0s - loss: 153.4932 - val_loss: 150.6499
Epoch 36/50
- 0s - loss: 151.3864 - val_loss: 147.8129
Epoch 37/50
- 0s - loss: 148.0773 - val_loss: 145.9286
Epoch 38/50
- 0s - loss: 145.7805 - val_loss: 142.1397
Epoch 39/50
- 0s - loss: 143.3021 - val_loss: 137.9944
Epoch 40/50
- 0s - loss: 140.0806 - val_loss: 138.1581
Epoch 41/50
- 0s - loss: 137.5236 - val_loss: 135.3461
Epoch 42/50
- 0s - loss: 134.8125 - val_loss: 133.8386
Epoch 43/50
- 0s - loss: 132.1174 - val_loss: 130.8983
Epoch 44/50
- 0s - loss: 129.3308 - val_loss: 127.3620
Epoch 45/50
- 0s - loss: 126.1189 - val_loss: 126.8274
Epoch 46/50
- 0s - loss: 123.1578 - val_loss: 125.4609
Epoch 47/50
- 0s - loss: 120.1846 - val_loss: 123.1138
Epoch 48/50
- 0s - loss: 117.5718 - val_loss: 122.3964
Epoch 49/50
- 0s - loss: 114.9057 - val_loss: 118.5039
Epoch 50/50
- 0s - loss: 111.8043 - val_loss: 118.2649
```

Out[10]:

```
<keras.callbacks.History at 0x7f2451ab4828>
```

You can refer to this [link](<https://keras.io/models/sequential/>) to learn about other functions that you can use for prediction or evaluation.

In [11]:

```
model = regression_model()
```


In [12]:

```
model.fit(predictors_norm, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

- 1s - loss: 1698.2179 - val_loss: 1215.7188

Epoch 2/50

- 0s - loss: 1663.4352 - val_loss: 1190.7206

Epoch 3/50

- 0s - loss: 1615.9694 - val_loss: 1150.7228

Epoch 4/50

- 0s - loss: 1540.8572 - val_loss: 1083.7393

Epoch 5/50

- 0s - loss: 1421.3596 - val_loss: 983.2373

Epoch 6/50

- 0s - loss: 1248.8737 - val_loss: 846.3096

Epoch 7/50

- 0s - loss: 1018.6126 - val_loss: 678.3351

Epoch 8/50

- 0s - loss: 755.0307 - val_loss: 496.8488

Epoch 9/50

- 0s - loss: 509.7421 - val_loss: 342.3384

Epoch 10/50

- 0s - loss: 346.9056 - val_loss: 244.8836

Epoch 11/50

- 0s - loss: 272.3920 - val_loss: 204.4806

Epoch 12/50

- 0s - loss: 249.0687 - val_loss: 188.1479

Epoch 13/50

- 0s - loss: 236.2585 - val_loss: 182.8122

Epoch 14/50

- 0s - loss: 226.3362 - val_loss: 177.2853

Epoch 15/50

- 0s - loss: 217.6977 - val_loss: 175.5469

Epoch 16/50

- 0s - loss: 210.9169 - val_loss: 171.9349

Epoch 17/50

- 0s - loss: 204.5435 - val_loss: 168.7015

Epoch 18/50

- 0s - loss: 199.2226 - val_loss: 166.5442

Epoch 19/50

- 0s - loss: 194.1455 - val_loss: 166.6936

Epoch 20/50

- 0s - loss: 190.3973 - val_loss: 164.6761

Epoch 21/50

- 0s - loss: 186.4633 - val_loss: 165.6336

Epoch 22/50

- 0s - loss: 183.4405 - val_loss: 164.7956

Epoch 23/50

- 0s - loss: 180.0897 - val_loss: 163.7727

Epoch 24/50

- 0s - loss: 177.5580 - val_loss: 163.8389

Epoch 25/50

- 0s - loss: 174.8561 - val_loss: 163.9029

Epoch 26/50

- 0s - loss: 172.6467 - val_loss: 162.2881

Epoch 27/50

- 0s - loss: 170.4997 - val_loss: 163.4355

Epoch 28/50

- 0s - loss: 168.3390 - val_loss: 163.0648

```
Epoch 29/50
- 0s - loss: 166.4030 - val_loss: 161.0971
Epoch 30/50
- 0s - loss: 165.2311 - val_loss: 163.0435
Epoch 31/50
- 0s - loss: 163.0253 - val_loss: 162.8007
Epoch 32/50
- 0s - loss: 161.4569 - val_loss: 163.9248
Epoch 33/50
- 0s - loss: 160.0337 - val_loss: 163.0733
Epoch 34/50
- 0s - loss: 158.2910 - val_loss: 164.2659
Epoch 35/50
- 0s - loss: 156.9108 - val_loss: 164.2933
Epoch 36/50
- 0s - loss: 155.6978 - val_loss: 164.6828
Epoch 37/50
- 0s - loss: 154.7859 - val_loss: 165.0226
Epoch 38/50
- 0s - loss: 153.0035 - val_loss: 166.9518
Epoch 39/50
- 0s - loss: 151.9325 - val_loss: 166.4964
Epoch 40/50
- 0s - loss: 151.0522 - val_loss: 167.9943
Epoch 41/50
- 0s - loss: 149.5892 - val_loss: 166.5662
Epoch 42/50
- 0s - loss: 148.5128 - val_loss: 168.2066
Epoch 43/50
- 0s - loss: 147.6377 - val_loss: 169.7690
Epoch 44/50
- 0s - loss: 146.5488 - val_loss: 170.3760
Epoch 45/50
- 0s - loss: 145.7479 - val_loss: 168.8901
Epoch 46/50
- 0s - loss: 145.3422 - val_loss: 171.4515
Epoch 47/50
- 0s - loss: 144.0320 - val_loss: 171.3076
Epoch 48/50
- 0s - loss: 142.9652 - val_loss: 173.2709
Epoch 49/50
- 0s - loss: 142.1486 - val_loss: 173.2225
Epoch 50/50
- 0s - loss: 141.5812 - val_loss: 173.2362
```

Out[12]:

<keras.callbacks.History at 0x7f245037ed30>

In [13]:

```
model = regression_model()
```

In [14]:

```
model.fit(predictors_norm, target, validation_split=0.3, epochs=50, verbose=2)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

- 1s - loss: 1704.5568 - val_loss: 1231.4729

Epoch 2/50

- 0s - loss: 1687.8204 - val_loss: 1221.6282

Epoch 3/50

- 0s - loss: 1668.1951 - val_loss: 1208.6375

Epoch 4/50

- 0s - loss: 1640.5399 - val_loss: 1189.0441

Epoch 5/50

- 0s - loss: 1599.9645 - val_loss: 1159.2489

Epoch 6/50

- 0s - loss: 1537.8771 - val_loss: 1113.5315

Epoch 7/50

- 0s - loss: 1449.6832 - val_loss: 1050.2998

Epoch 8/50

- 0s - loss: 1330.4825 - val_loss: 965.3668

Epoch 9/50

- 0s - loss: 1172.7266 - val_loss: 858.7152

Epoch 10/50

- 0s - loss: 982.5863 - val_loss: 725.2443

Epoch 11/50

- 0s - loss: 765.9098 - val_loss: 576.4156

Epoch 12/50

- 0s - loss: 553.9582 - val_loss: 426.7120

Epoch 13/50

- 0s - loss: 382.7408 - val_loss: 308.7035

Epoch 14/50

- 0s - loss: 282.0862 - val_loss: 237.1588

Epoch 15/50

- 0s - loss: 243.7353 - val_loss: 202.8895

Epoch 16/50

- 0s - loss: 223.9785 - val_loss: 193.2828

Epoch 17/50

- 0s - loss: 212.7122 - val_loss: 184.7323

Epoch 18/50

- 0s - loss: 203.3505 - val_loss: 176.4084

Epoch 19/50

- 0s - loss: 197.0796 - val_loss: 171.1934

Epoch 20/50

- 0s - loss: 191.5619 - val_loss: 165.3285

Epoch 21/50

- 0s - loss: 187.4990 - val_loss: 157.9938

Epoch 22/50

- 0s - loss: 182.6198 - val_loss: 157.2669

Epoch 23/50

- 0s - loss: 179.4335 - val_loss: 154.8402

Epoch 24/50

- 0s - loss: 176.6428 - val_loss: 151.2046

Epoch 25/50

- 0s - loss: 174.2190 - val_loss: 146.9907

Epoch 26/50

- 0s - loss: 172.0937 - val_loss: 145.6086

Epoch 27/50

- 0s - loss: 169.8503 - val_loss: 143.5039

Epoch 28/50

- 0s - loss: 167.7833 - val_loss: 142.0920

```
Epoch 29/50
- 0s - loss: 165.9921 - val_loss: 141.5505
Epoch 30/50
- 0s - loss: 164.0796 - val_loss: 140.6397
Epoch 31/50
- 0s - loss: 162.3230 - val_loss: 138.0121
Epoch 32/50
- 0s - loss: 160.7015 - val_loss: 137.8299
Epoch 33/50
- 0s - loss: 158.6024 - val_loss: 136.5406
Epoch 34/50
- 0s - loss: 156.2944 - val_loss: 136.4598
Epoch 35/50
- 0s - loss: 154.2929 - val_loss: 135.3828
Epoch 36/50
- 0s - loss: 152.3761 - val_loss: 136.0370
Epoch 37/50
- 0s - loss: 150.0692 - val_loss: 137.7441
Epoch 38/50
- 0s - loss: 148.0991 - val_loss: 137.8791
Epoch 39/50
- 0s - loss: 145.7735 - val_loss: 137.3048
Epoch 40/50
- 0s - loss: 143.6673 - val_loss: 138.0521
Epoch 41/50
- 0s - loss: 141.8274 - val_loss: 138.6647
Epoch 42/50
- 0s - loss: 139.7745 - val_loss: 139.2951
Epoch 43/50
- 0s - loss: 137.7955 - val_loss: 140.7642
Epoch 44/50
- 0s - loss: 136.0377 - val_loss: 141.7609
Epoch 45/50
- 1s - loss: 134.4538 - val_loss: 140.9436
Epoch 46/50
- 0s - loss: 132.6085 - val_loss: 144.4127
Epoch 47/50
- 0s - loss: 131.1760 - val_loss: 143.7522
Epoch 48/50
- 0s - loss: 129.4891 - val_loss: 147.4993
Epoch 49/50
- 0s - loss: 127.9810 - val_loss: 146.4812
Epoch 50/50
- 0s - loss: 126.4801 - val_loss: 148.7281
```

Out[14]:

```
<keras.callbacks.History at 0x7f24486a26d8>
```

In [15]:

```
score = model.evaluate(predictors_norm, target)
```

```
1030/1030 [=====] - 0s 118us/step
```

Feel free to vary the following and note what impact each change has on the model's performance:

1. Increase or decrease number of neurons in hidden layers
2. Add more hidden layers
3. Increase number of epochs

In [16]:

```
target.mean()
```

Out[16]:

```
35.817961165048544
```

In [17]:

```
target.std()
```

Out[17]:

```
16.705741961912512
```

In [19]:

```
predictors_norm.mean()
```

Out[19]:

Cement	2.432224e-15
Blast Furnace Slag	-8.513686e-16
Fly Ash	3.837815e-16
Water	1.846743e-15
Superplasticizer	-9.641155e-16
Coarse Aggregate	6.818710e-15
Fine Aggregate	1.232571e-14
Age	3.640022e-16

dtype: float64

In [20]:

```
predictors_norm.std()
```

Out[20]:

Cement	1.0
Blast Furnace Slag	1.0
Fly Ash	1.0
Water	1.0
Superplasticizer	1.0
Coarse Aggregate	1.0
Fine Aggregate	1.0
Age	1.0

dtype: float64

In []: