



Watson Studio democratizes machine learning and deep learning to accelerate infusion of AI in your business to drive innovation. Watson Studio provides a suite of tools and a collaborative environment for data scientists, developers and domain experts.

(http://cocl.us/pytorch_link_top)



Linear regression 1D: Training Two Parameter Stochastic Gradient Descent (SGD)

Table of Contents

In this Lab, you will practice training a model by using Stochastic Gradient descent.

- [Make Some Data](#)
- [Create the Model and Cost Function \(Total Loss\)](#)
- [Train the Model:Batch Gradient Descent](#)
- [Train the Model:Stochastic gradient descent](#)
- [Train the Model:Stochastic gradient descent with Data Loader](#)

Estimated Time Needed: **30 min**

Preparation

We'll need the following libraries:

In [1]:

```
# These are the libraries we are going to use in the lab.

import torch
import matplotlib.pyplot as plt
import numpy as np

from mpl_toolkits import mplot3d
```

The class `plot_error_surfaces` is just to help you visualize the data space and the parameter space during training and has nothing to do with PyTorch.

In [2]:

```
# The class for plot the diagram
```

```
class plot_error_surfaces(object):
```

```
    # Constructor
```

```
    def __init__(self, w_range, b_range, X, Y, n_samples = 30, go = True):
```

```
        W = np.linspace(-w_range, w_range, n_samples)
```

```
        B = np.linspace(-b_range, b_range, n_samples)
```

```
        w, b = np.meshgrid(W, B)
```

```
        Z = np.zeros((30, 30))
```

```
        count1 = 0
```

```
        self.y = Y.numpy()
```

```
        self.x = X.numpy()
```

```
        for w1, b1 in zip(w, b):
```

```
            count2 = 0
```

```
            for w2, b2 in zip(w1, b1):
```

```
                Z[count1, count2] = np.mean((self.y - w2 * self.x + b2) ** 2)
```

```
                count2 += 1
```

```
            count1 += 1
```

```
        self.Z = Z
```

```
        self.w = w
```

```
        self.b = b
```

```
        self.W = []
```

```
        self.B = []
```

```
        self.LOSS = []
```

```
        self.n = 0
```

```
        if go == True:
```

```
            plt.figure()
```

```
            plt.figure(figsize = (7.5, 5))
```

```
            plt.axes(projection = '3d').plot_surface(self.w, self.b, self.Z, rstride
```

```
e = 1, cstride = 1, cmap = 'viridis', edgecolor = 'none')
```

```
            plt.title('Loss Surface')
```

```
            plt.xlabel('w')
```

```
            plt.ylabel('b')
```

```
            plt.show()
```

```
            plt.figure()
```

```
            plt.title('Loss Surface Contour')
```

```
            plt.xlabel('w')
```

```
            plt.ylabel('b')
```

```
            plt.contour(self.w, self.b, self.Z)
```

```
            plt.show()
```

```
    # Setter
```

```
    def set_para_loss(self, W, B, loss):
```

```
        self.n = self.n + 1
```

```
        self.W.append(W)
```

```
        self.B.append(B)
```

```
        self.LOSS.append(loss)
```

```
    # Plot diagram
```

```
    def final_plot(self):
```

```
        ax = plt.axes(projection = '3d')
```

```
        ax.plot_wireframe(self.w, self.b, self.Z)
```

```
        ax.scatter(self.W, self.B, self.LOSS, c = 'r', marker = 'x', s = 200, alpha
```

```
= 1)
```

```

plt.figure()
plt.contour(self.w, self.b, self.Z)
plt.scatter(self.W, self.B, c = 'r', marker = 'x')
plt.xlabel('w')
plt.ylabel('b')
plt.show()

# Plot diagram
def plot_ps(self):
    plt.subplot(121)
    plt.ylim
    plt.plot(self.x, self.y, 'ro', label = "training points")
    plt.plot(self.x, self.W[-1] * self.x + self.B[-1], label = "estimated line"
)

    plt.xlabel('x')
    plt.ylabel('y')
    plt.ylim((-10, 15))
    plt.title('Data Space Iteration: ' + str(self.n))
    plt.subplot(122)
    plt.contour(self.w, self.b, self.Z)
    plt.scatter(self.W, self.B, c = 'r', marker = 'x')
    plt.title('Loss Surface Contour Iteration' + str(self.n))
    plt.xlabel('w')
    plt.ylabel('b')
    plt.show()

```

Make Some Data

Set random seed:

In [3]:

```

# Set random seed

torch.manual_seed(1)

```

Out[3]:

```

<torch._C.Generator at 0x7fa389e482f0>

```

Generate values from -3 to 3 that create a line with a slope of 1 and a bias of -1. This is the line that you need to estimate. Add some noise to the data:

In [4]:

```

# Setup the actual data and simulated data

X = torch.arange(-3, 3, 0.1).view(-1, 1)
f = 1 * X - 1
Y = f + 0.1 * torch.randn(X.size())


```

Plot the results:

In [5]:

```
# Plot out the data dots and line

plt.plot(X.numpy(), Y.numpy(), 'rx', label = 'y')
plt.plot(X.numpy(), f.numpy(), label = 'f')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
plt.show()
```



Create the Model and Cost Function (Total Loss)

Define the forward function:

In [6]:

```
# Define the forward function

def forward(x):
    return w * x + b
```

Define the cost or criterion function (MSE):

In [7]:

```
# Define the MSE Loss function



def criterion(yhat, y):
    return torch.mean((yhat - y) ** 2)
```

Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

In [8]:

```
# Create plot_error_surfaces for viewing the data  
get_surface = plot_error_surfaces(15, 13, X, Y, 30)
```

<Figure size 432x288 with 0 Axes>



Train the Model: Batch Gradient Descent

Create model parameters w , b by setting the argument `requires_grad` to True because the system must learn it.

In [9]:

```
# Define the parameters w, b for y = wx + b

w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)
```

Set the learning rate to 0.1 and create an empty list `LOSS` for storing the loss for each iteration.

In [10]:

```
# Define learning rate and create an empty list for containing the loss for each iteration.

lr = 0.1
LOSS_BGD = []
```

Define `train_model` function for train the model.

In [11]:

```
# The function for training the model

def train_model(iter):

    # Loop
    for epoch in range(iter):

        # make a prediction
        Yhat = forward(X)

        # calculate the loss
        loss = criterion(Yhat, Y)

        # Section for plotting
        get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist())
        get_surface.plot_ps()

        # store the loss in the list LOSS_BGD
        LOSS_BGD.append(loss)

        # backward pass: compute gradient of the loss with respect to all the learnable parameters
        loss.backward()

        # update parameters slope and bias
        w.data = w.data - lr * w.grad.data
        b.data = b.data - lr * b.grad.data

        # zero the gradients before running the backward pass
        w.grad.data.zero_()
        b.grad.data.zero_()
```


Run 10 epochs of batch gradient descent: **bug** data space is 1 iteration ahead of parameter space.

In [12]:


```
# Train the model with 10 iterations
```

```
train_model(10)
```


Data Space Iteration: 1




Loss Surface Contour Iteration1





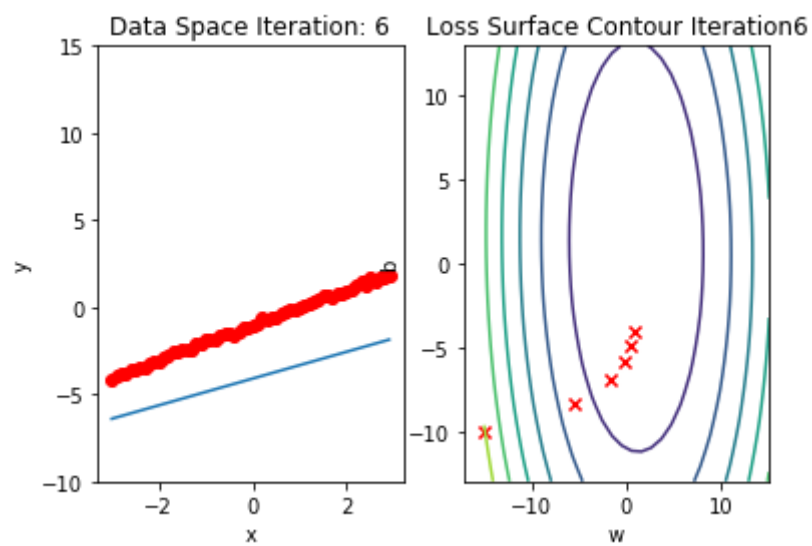

Data Space Iteration: 2







Loss Surface Contour Iteration2











Train the Model: Stochastic Gradient Descent

Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

In [13]:

```
# Create plot_error_surfaces for viewing the data
get_surface = plot_error_surfaces(15, 13, X, Y, 30, go = False)
```

Define `train_model_SGD` function for training the model.

In [14]:

```
# The function for training the model

LOSS_SGD = []
w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)

def train_model_SGD(iter):

    # Loop
    for epoch in range(iter):

        # SGD is an approximation of our true total loss/cost, in this line of code
        we calculate our true loss/cost and store it
        Yhat = forward(X)

        # store the loss
        LOSS_SGD.append(criterion(Yhat, Y).tolist())

        for x, y in zip(X, Y):

            # make a prediction
            yhat = forward(x)

            # calculate the loss
            loss = criterion(yhat, y)

            # Section for plotting
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist()
            ())

            # backward pass: compute gradient of the loss with respect to all the learnable parameters
            loss.backward()

            # update parameters slope and bias
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr * b.grad.data

            # zero the gradients before running the backward pass
            w.grad.data.zero_()
            b.grad.data.zero_()

        #plot surface and data space after each epoch
        get_surface.plot_ps()
```


Run 10 epochs of stochastic gradient descent: **bug** data space is 1 iteration ahead of parameter space.

In [15]:


```
# Train the model with 10 iterations
```

```
train_model_SGD(10)
```



Data Space Iteration: 60 Loss Surface Contour Iteration60




Data Space Iteration: 120 Loss Surface Contour Iteration120




Data Space Iteration: 180 Loss Surface Contour Iteration180




Data Space Iteration: 240 Loss Surface Contour Iteration240




Data Space Iteration: 300 Loss Surface Contour Iteration300




Data Space Iteration: 360 Loss Surface Contour Iteration360




Data Space Iteration: 420 Loss Surface Contour Iteration420




Data Space Iteration: 480 Loss Surface Contour Iteration480



Data Space Iteration: 540 Loss Surface Contour Iteration540






Compare the loss of both batch gradient descent as SGD.

In [16]:

```
# Plot out the LOSS_BGD and LOSS_SGD

plt.plot(LOSS_BGD, label = "Batch Gradient Descent")
plt.plot(LOSS_SGD, label = "Stochastic Gradient Descent")
plt.xlabel('epoch')
plt.ylabel('Cost/ total loss')
plt.legend()
plt.show()
```



SGD with Dataset DataLoader

Import the module for building a dataset class:

In [17]:

```
# Import the library for DataLoader  
  
from torch.utils.data import Dataset, DataLoader
```

Create a dataset class:

In [18]:

```
# Dataset Class  
  
class Data(Dataset):  
  
    # Constructor  
    def __init__(self):  
        self.x = torch.arange(-3, 3, 0.1).view(-1, 1)  
        self.y = 1 * self.x - 1  
        self.len = self.x.shape[0]  
  
    # Getter  
    def __getitem__(self, index):  
        return self.x[index], self.y[index]  
  
    # Return the length  
    def __len__(self):  
        return self.len
```

Create a dataset object and check the length of the dataset.

In [19]:

```
# Create the dataset and check the length  
  
dataset = Data()  
print("The length of dataset: ", len(dataset))
```

The length of dataset: 60

Obtain the first training point:

In [20]:

```
# Print the first point

x, y = dataset[0]
print("(" , x, " , " , y, ")")

( tensor([-3.]) , tensor([-4.]) )
```

Similarly, obtain the first three training points:

In [21]:

```
# Print the first 3 point

x, y = dataset[0:3]
print("The first 3 x: " , x)
print("The first 3 y: " , y)
```

```
The first 3 x: tensor([[ -3.0000],
                    [-2.9000],
                    [-2.8000]])
The first 3 y: tensor([[ -4.0000],
                    [-3.9000],
                    [-3.8000]])
```

Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

In [22]:

```
# Create plot_error_surfaces for viewing the data

get_surface = plot_error_surfaces(15, 13, X, Y, 30, go = False)
```

Create a `DataLoader` object by using the constructor:

In [23]:

```
# Create DataLoader

trainloader = DataLoader(dataset = dataset, batch_size = 1)
```

Define `train_model_DataLoader` function for training the model.

In [24]:

```
# The function for training the model

w = torch.tensor(-15.0,requires_grad=True)
b = torch.tensor(-10.0,requires_grad=True)
LOSS_Loader = []

def train_model_DataLoader(epochs):

    # Loop
    for epoch in range(epochs):

        # SGD is an approximation of our true total loss/cost, in this line of code
we calculate our true loss/cost and store it
        Yhat = forward(X)

        # store the loss
        LOSS_Loader.append(criterion(Yhat, Y).tolist())

        for x, y in trainloader:

            # make a prediction
            yhat = forward(x)

            # calculate the loss
            loss = criterion(yhat, y)

            # Section for plotting
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist
            ())

            # Backward pass: compute gradient of the loss with respect to all the l
earnable parameters
            loss.backward()

            # Update parameters slope
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr* b.grad.data

            # Clear gradients
            w.grad.data.zero_()
            b.grad.data.zero_()

        #plot surface and data space after each epoch
        get_surface.plot_ps()
```


Run 10 epochs of stochastic gradient descent: **bug** data space is 1 iteration ahead of parameter space.

In [25]:


```
# Run 10 iterations
```

```
train_model_DataLoader(10)
```



Data Space Iteration: 60 Loss Surface Contour Iteration60




Data Space Iteration: 120 Loss Surface Contour Iteration120




Data Space Iteration: 180 Loss Surface Contour Iteration180




Data Space Iteration: 240 Loss Surface Contour Iteration240




Data Space Iteration: 300 Loss Surface Contour Iteration300




Data Space Iteration: 360 Loss Surface Contour Iteration360




Data Space Iteration: 420 Loss Surface Contour Iteration420




Data Space Iteration: 480 Loss Surface Contour Iteration480



Data Space Iteration: 540 Loss Surface Contour Iteration540






Compare the loss of both batch gradient decent as SGD. Note that SGD converges to a minimum faster, that is, it decreases faster.

In [26]:

```
# Plot the LOSS_BGD and LOSS_Loader
```

```
plt.plot(LOSS_BGD,label="Batch Gradient Descent")
plt.plot(LOSS_Loader,label="Stochastic Gradient Descent with DataLoader")
plt.xlabel('epoch')
plt.ylabel('Cost/ total loss')
plt.legend()
plt.show()
```



Practice

For practice, try to use SGD with DataLoader to train model with 10 iterations. Store the total loss in `loss`. We are going to use it in the next question.


In [27]:

```
# Practice: Use SGD with trainloader to train model and store the total loss in LOSS
```


```
LOSS = []
w = torch.tensor(-12.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)

def my_train_model(epochs):
    for epoch in range(epochs):
        Yhat = forward(X)
        LOSS.append(criterion(Yhat, X))
        for x, y in trainloader:
            yhat = forward(x)
            loss = criterion(yhat, y)
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist())
        loss.backward()
        w.data = w.data - lr * w.grad.data
        b.data = b.data - lr * b.grad.data
        w.grad.data.zero_()
        b.grad.data.zero_()
        get_surface.plot_ps()
my_train_model(10)
```


Data Space Iteration: 660 Loss Surface Contour Iteration660




Data Space Iteration: 720 Loss Surface Contour Iteration720




Data Space Iteration: 780 Loss Surface Contour Iteration780




Data Space Iteration: 840 Loss Surface Contour Iteration840




Data Space Iteration: 900 Loss Surface Contour Iteration900




Data Space Iteration: 960 Loss Surface Contour Iteration960





Data Space Iteration: 1020 Loss Surface Contour Iteration 1020



Data Space Iteration: 1080 Loss Surface Contour Iteration 1080






Double-click **here** for the solution.

Plot the total loss

In [28]:

```
# Practice: Plot the total loss using LOSS

plt.plot(LOSS, label = 'Stochastic Gradient Descent')
plt.xlabel('iteration')
plt.ylabel('Cost/total loss')
plt.legend()
plt.show()
```



Double-click **here** for the solution.

Get IBM Watson Studio free of charge!

Build and train AI & machine learning models, prepare and analyze data – all in a flexible, hybrid cloud environment. Get IBM Watson Studio Lite Plan free of charge.



Learn

Get started or get better with built-in learning.



Create

Use the best of open source tooling with IBM innovation.



Collaborate

Work smarter using community, work faster with your team.

[Sign Up For a Free Trial](#)

(<http://cocl.us/pytorch> link bottom)

About the Authors:

[Joseph Santarcangelo](https://www.linkedin.com/in/joseph-s-50398b136/) (<https://www.linkedin.com/in/joseph-s-50398b136/>) has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Other contributors: [Michelle Carey](https://www.linkedin.com/in/michelleccarey/) (<https://www.linkedin.com/in/michelleccarey/>), [Mavis Zhou](https://www.linkedin.com/in/jiahui-mavis-zhou-a4537814a) (www.linkedin.com/in/jiahui-mavis-zhou-a4537814a).

Thanks to: Andrew Kin ,Alessandro Barboza

Copyright © 2018 cognitiveclass.ai (cognitiveclass.ai?utm_source=bducopyrightlink&utm_medium=dswb&utm_campaign=bdu). This notebook and its source code are released under the terms of the [MIT License](https://bigdatauniversity.com/mit-license/) (<https://bigdatauniversity.com/mit-license/>).

In []: