



Watson Studio democratizes machine learning and deep learning to accelerate infusion of AI in your business to drive innovation. Watson Studio provides a suite of tools and a collaborative environment for data scientists, developers and domain experts.

(http://cocl.us/pytorch_link_top)



Linear Regression 1D: Training Two Parameter Mini-Batch Gradient Decent

Table of Contents

In this Lab, you will practice training a model by using Mini-Batch Gradient Descent.

- [Make Some Data](#)
- [Create the Model and Cost Function \(Total Loss\)](#)
- [Train the Model: Batch Gradient Descent](#)
- [Train the Model: Stochastic Gradient Descent with Dataset DataLoader](#)
- [Train the Model: Mini Batch Gradient Decent: Batch Size Equals 5](#)
- [Train the Model: Mini Batch Gradient Decent: Batch Size Equals 10](#)

Estimated Time Needed: **30 min**

</div>

Preparation

We'll need the following libraries:

In [1]:

```
# Import the libraries we need for this lab

import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits import mplot3d
```

The class `plot_error_surfaces` is just to help you visualize the data space and the parameter space during training and has nothing to do with PyTorch.

In [2]:

```
# The class for plotting the diagrams
```

```
class plot_error_surfaces(object):
```

```
    # Constructor
```

```
    def __init__(self, w_range, b_range, X, Y, n_samples = 30, go = True):
```

```
        W = np.linspace(-w_range, w_range, n_samples)
```

```
        B = np.linspace(-b_range, b_range, n_samples)
```

```
        w, b = np.meshgrid(W, B)
```

```
        Z = np.zeros((30, 30))
```

```
        count1 = 0
```

```
        self.y = Y.numpy()
```

```
        self.x = X.numpy()
```

```
        for w1, b1 in zip(w, b):
```

```
            count2 = 0
```

```
            for w2, b2 in zip(w1, b1):
```

```
                Z[count1, count2] = np.mean((self.y - w2 * self.x + b2) ** 2)
```

```
                count2 += 1
```

```
            count1 += 1
```

```
        self.Z = Z
```

```
        self.w = w
```

```
        self.b = b
```

```
        self.W = []
```

```
        self.B = []
```

```
        self.LOSS = []
```

```
        self.n = 0
```

```
        if go == True:
```

```
            plt.figure()
```

```
            plt.figure(figsize = (7.5, 5))
```

```
            plt.axes(projection = '3d').plot_surface(self.w, self.b, self.Z, rstride
```

```
e = 1, cstride = 1, cmap = 'viridis', edgecolor = 'none')
```

```
            plt.title('Loss Surface')
```

```
            plt.xlabel('w')
```

```
            plt.ylabel('b')
```

```
            plt.show()
```

```
            plt.figure()
```

```
            plt.title('Loss Surface Contour')
```

```
            plt.xlabel('w')
```

```
            plt.ylabel('b')
```

```
            plt.contour(self.w, self.b, self.Z)
```

```
            plt.show()
```

```
    # Setter
```

```
    def set_para_loss(self, W, B, loss):
```

```
        self.n = self.n + 1
```

```
        self.W.append(W)
```

```
        self.B.append(B)
```

```
        self.LOSS.append(loss)
```

```
    # Plot diagram
```

```
    def final_plot(self):
```

```
        ax = plt.axes(projection = '3d')
```

```
        ax.plot_wireframe(self.w, self.b, self.Z)
```

```
        ax.scatter(self.W, self.B, self.LOSS, c = 'r', marker = 'x', s = 200, alpha
```

```
= 1)
```

```

plt.figure()
plt.contour(self.w, self.b, self.Z)
plt.scatter(self.W, self.B, c = 'r', marker = 'x')
plt.xlabel('w')
plt.ylabel('b')
plt.show()

# Plot diagram
def plot_ps(self):
    plt.subplot(121)
    plt.ylim()
    plt.plot(self.x, self.y, 'ro', label = "training points")
    plt.plot(self.x, self.W[-1] * self.x + self.B[-1], label = "estimated line"
)

    plt.xlabel('x')
    plt.ylabel('y')
    plt.title('Data Space Iteration: ' + str(self.n))
    plt.subplot(122)
    plt.contour(self.w, self.b, self.Z)
    plt.scatter(self.W, self.B, c = 'r', marker = 'x')
    plt.title('Loss Surface Contour')
    plt.xlabel('w')
    plt.ylabel('b')
    plt.show()

```

Make Some Data

Import PyTorch and set random seed:

In [3]:

```

# Import PyTorch library

import torch
torch.manual_seed(1)

```

Out[3]:

```
<torch._C.Generator at 0x7f23772dc050>
```

Generate values from -3 to 3 that create a line with a slope of 1 and a bias of -1. This is the line that you need to estimate. Add some noise to the data:

In [4]:

```

# Generate the data with noise and the line

X = torch.arange(-3, 3, 0.1).view(-1, 1)
f = 1 * X - 1
Y = f + 0.1 * torch.randn(X.size())


```

Plot the results:

In [5]:

```
# Plot the line and the data

plt.plot(X.numpy(), Y.numpy(), 'rx', label = 'y')
plt.plot(X.numpy(), f.numpy(), label = 'f')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
plt.show()
```



Create the Model and Cost Function (Total Loss)

Define the forward function:

In [6]:

```
# Define the prediction function

def forward(x):
    return w * x + b
```

Define the cost or criterion function:

In [7]:

```
# Define the cost function

def criterion(yhat, y):
    return torch.mean((yhat - y) ** 2)
```



Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

In [8]:

```
# Create a plot_error_surfaces object.
```

```
get_surface = plot_error_surfaces(15, 13, X, Y, 30)
```

<Figure size 432x288 with 0 Axes>



Train the Model: Batch Gradient Descent (BGD)

Define `train_model_BGD` function.

In [9]:

```
# Define the function for training model

w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)
lr = 0.1
LOSS_BGD = []

def train_model_BGD(epochs):
    for epoch in range(epochs):
        Yhat = forward(X)
        loss = criterion(Yhat, Y)
        LOSS_BGD.append(loss)
        get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist())
        get_surface.plot_ps()
        loss.backward()
        w.data = w.data - lr * w.grad.data
        b.data = b.data - lr * b.grad.data
        w.grad.data.zero_()
        b.grad.data.zero_()
```


Run 10 epochs of batch gradient descent: **bug** data space is 1 iteration ahead of parameter space.

In [10]:


```
# Run train_model_BGD with 10 iterations
```

```
train_model_BGD(10)
```



Data Space Iteration: 1




Loss Surface Contour




Data Space Iteration: 2




Loss Surface Contour




Data Space Iteration: 3




Loss Surface Contour




Data Space Iteration: 4




Loss Surface Contour




Data Space Iteration: 5




Loss Surface Contour




Data Space Iteration: 6




Loss Surface Contour




Data Space Iteration: 7




Loss Surface Contour

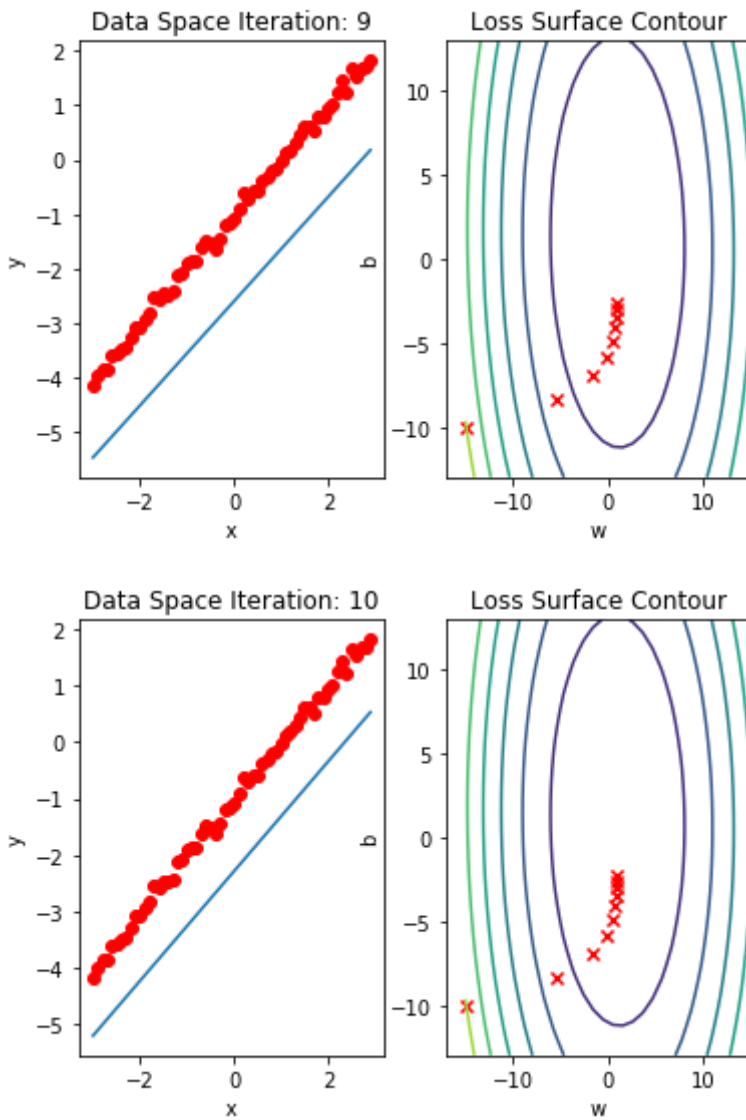


Data Space Iteration: 8



Loss Surface Contour





Stochastic Gradient Descent (SGD) with Dataset DataLoader

Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

In [11]:

```
# Create a plot_error_surfaces object.  
get_surface = plot_error_surfaces(15, 13, X, Y, 30, go = False)
```

Import Dataset and DataLoader libraries

In [12]:

```
# Import libraries

from torch.utils.data import Dataset, DataLoader
```

Create Data class

In [13]:

```
# Create class Data

class Data(Dataset):

    # Constructor
    def __init__(self):
        self.x = torch.arange(-3, 3, 0.1).view(-1, 1)
        self.y = 1 * X - 1
        self.len = self.x.shape[0]

    # Getter
    def __getitem__(self, index):
        return self.x[index], self.y[index]

    # Get length
    def __len__(self):
        return self.len
```

Create a dataset object and a dataloader object:

In [14]:

```
# Create Data object and DataLoader object

dataset = Data()
trainloader = DataLoader(dataset = dataset, batch_size = 1)
```

Define train_model_SGD function for training the model.

In [15]:

```
# Define train_model_SGD function

w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)
LOSS_SGD = []
lr = 0.1
def train_model_SGD(epochs):
    for epoch in range(epochs):
        Yhat = forward(X)
        get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), criterion(Yhat,
Y).tolist())
        get_surface.plot_ps()
        LOSS_SGD.append(criterion(forward(X), Y).tolist())
        for x, y in trainloader:
            yhat = forward(x)
            loss = criterion(yhat, y)
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist
())


            loss.backward()
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr * b.grad.data
            w.grad.data.zero_()
            b.grad.data.zero_()
        get_surface.plot_ps()
```

Run 10 epochs of stochastic gradient descent: **bug** data space is 1 iteration ahead of parameter space.


In [16]:

```
# Run train_model_SGD(iter) with 10 iterations  
train_model_SGD(10)
```



Data Space Iteration: 1




Loss Surface Contour




Data Space Iteration: 61




Loss Surface Contour




Data Space Iteration: 62




Loss Surface Contour




Data Space Iteration: 122




Loss Surface Contour




Data Space Iteration: 123




Loss Surface Contour




Data Space Iteration: 183




Loss Surface Contour




Data Space Iteration: 184




Loss Surface Contour




Data Space Iteration: 244




Loss Surface Contour






Data Space Iteration: 245




Loss Surface Contour







Data Space Iteration: 367




Loss Surface Contour




Data Space Iteration: 427




Loss Surface Contour






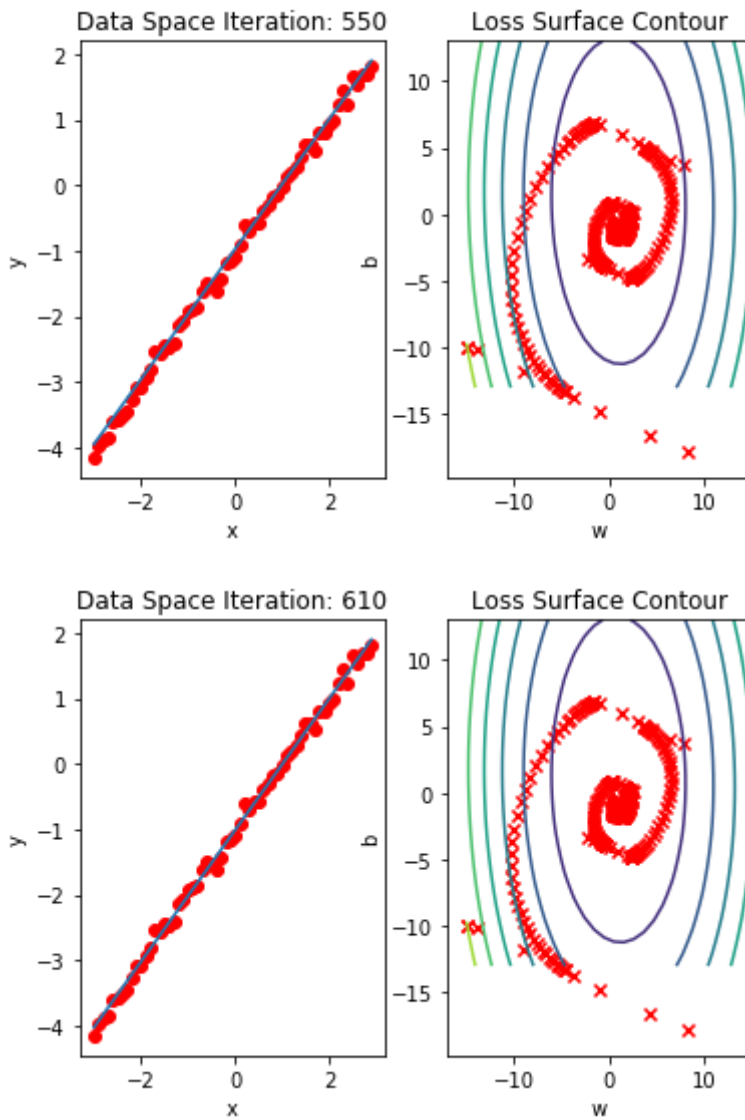
Data Space Iteration: 428



Loss Surface Contour







Mini Batch Gradient Descent: Batch Size Equals 5

Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

In [17]:

```
# Create a plot_error_surfaces object.

get_surface = plot_error_surfaces(15, 13, X, Y, 30, go = False)
```

Create `Data` object and create a `Dataloader` object where the batch size equals 5:

In [18]:

```
# Create DataLoader object and Data object

dataset = Data()
trainloader = DataLoader(dataset = dataset, batch_size = 5)
```

Define `train_model_Mini5` function to train the model.

In [19]:

```
# Define train_model_Mini5 function

w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)
LOSS_MINI5 = []
lr = 0.1

def train_model_Mini5(epochs):
    for epoch in range(epochs):
        Yhat = forward(X)
        get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), criterion(Yhat,
Y).tolist())
        get_surface.plot_ps()
        LOSS_MINI5.append(criterion(forward(X), Y).tolist())
        for x, y in trainloader:
            yhat = forward(x)
            loss = criterion(yhat, y)
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist
())

            loss.backward()
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr * b.grad.data
            w.grad.data.zero_()
            b.grad.data.zero_()
```


Run 10 epochs of mini-batch gradient descent: **bug** data space is 1 iteration ahead of parameter space.

In [20]:


```
# Run train_model_Mini5 with 10 iterations.
```

```
train_model_Mini5(10)
```


Data Space Iteration: 1




Loss Surface Contour




Data Space Iteration: 14




Loss Surface Contour




Data Space Iteration: 27




Loss Surface Contour




Data Space Iteration: 40




Loss Surface Contour








Data Space Iteration: 53



Loss Surface Contour







Mini Batch Gradient Descent: Batch Size Equals 10

Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

In [21]:

```
# Create a plot_error_surfaces object.
```

```
get_surface = plot_error_surfaces(15, 13, X, Y, 30, go = False)
```

Create Data object and create a DataLoader object batch size equals 10

In [22]:

```
# Create DataLoader object

dataset = Data()
trainloader = DataLoader(dataset = dataset, batch_size = 10)
```

Define train_model_Mini10 function for training the model.

In [23]:

```
# Define train_model_Mini5 function

w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)
LOSS_MINI10 = []
lr = 0.1

def train_model_Mini10(epochs):
    for epoch in range(epochs):
        Yhat = forward(X)
        get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), criterion(Yhat,
Y).tolist())
        get_surface.plot_ps()
        LOSS_MINI10.append(criterion(forward(X),Y).tolist())
        for x, y in trainloader:
            yhat = forward(x)
            loss = criterion(yhat, y)
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist
            ())
            loss.backward()
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr * b.grad.data
            w.grad.data.zero_()
            b.grad.data.zero_()
```


Run 10 epochs of mini-batch gradient descent: **bug** data space is 1 iteration ahead of parameter space.

In [24]:


```
# Run train_model_Mini5 with 10 iterations.
```

```
train_model_Mini10(10)
```


Data Space Iteration: 1




Loss Surface Contour




Data Space Iteration: 8




Loss Surface Contour




Data Space Iteration: 15




Loss Surface Contour




Data Space Iteration: 22




Loss Surface Contour








Data Space Iteration: 29



Loss Surface Contour







Plot the loss for each epoch:


In [25]:

```
# Plot out the LOSS for each method
```

```
plt.plot(LOSS_BGD,label = "Batch Gradient Descent")  
plt.plot(LOSS_SGD,label = "Stochastic Gradient Descent")  
plt.plot(LOSS_MINI5,label = "Mini-Batch Gradient Descent, Batch size: 5")  
plt.plot(LOSS_MINI10,label = "Mini-Batch Gradient Descent, Batch size: 10")  
plt.legend()
```

Out[25]:

<matplotlib.legend.Legend at 0x7f23703b8b38>



Practice

Perform mini batch gradient descent with a batch size of 20. Store the total loss for each epoch in the list LOSS20.

In [26]:


```
# Practice: Perform mini batch gradient descent with a batch size of 20.

dataset = Data()
trainloader = DataLoader(dataset = dataset, batch_size = 20)
w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)


LOSS_MINI20 = []
lr = 0.1

def my_train_model(epochs):
    for epoch in range(epochs):
        Yhat = forward(X)
        get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), criterion(Yhat,
Y).tolist())
        get_surface.plot_ps()
        LOSS_MINI20.append(criterion(forward(X), Y).tolist())
        for x, y in trainloader:
            yhat = forward(x)
            loss =criterion(yhat, y)
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.tolist
())
            loss.backward()
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr * b.grad.data
            w.grad.data.zero_()
            b.grad.data.zero_()
my_train_model(10)
```


Data Space Iteration: 71




Loss Surface Contour




Data Space Iteration: 75




Loss Surface Contour




Data Space Iteration: 79




Loss Surface Contour




Data Space Iteration: 83




Loss Surface Contour




Data Space Iteration: 87




Loss Surface Contour




Data Space Iteration: 91




Loss Surface Contour




Data Space Iteration: 95




Loss Surface Contour





Data Space Iteration: 99



Loss Surface Contour





Double-click [here](#) for the solution.


Plot a graph that shows the LOSS results for all the methods.

In [27]:

```
# Practice: Plot a graph to show all the LOSS functions
plt.plot(LOSS_BGD, label = 'Batch Gradient Descent')
plt.plot(LOSS_SGD, label = 'Stochastic Gradient Descent')
plt.plot(LOSS_MINI5, label = 'Mini-Batch Gradient Descent, Batch size: 5')
plt.plot(LOSS_MINI10, label = 'Mini-Batch Gradient Descent, Batch size :10')
plt.plot(LOSS_MINI20, label = 'Mini-Batch Gradient Descent, Batch size :20')
plt.legend()
```

Out[27]:


<matplotlib.legend.Legend at 0x7f237007e4a8>



Double-click **here** for the solution.


Get IBM Watson Studio free of charge!

Build and train AI & machine learning models, prepare and analyze data – all in a flexible, hybrid cloud environment. Get IBM Watson Studio Lite Plan free of charge.




Learn

Get started or get better with built-in learning.



Create

Use the best of open source tooling with IBM innovation.



Collaborate

Work smarter using community, work faster with your team.

[Sign Up For a Free Trial](#)

(<http://cocl.us/pytorch> link bottom)

About the Authors:

[Joseph Santarcangelo](https://www.linkedin.com/in/joseph-s-50398b136/) (<https://www.linkedin.com/in/joseph-s-50398b136/>) has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Other contributors: [Michelle Carey](https://www.linkedin.com/in/michelleccarey/) (<https://www.linkedin.com/in/michelleccarey/>), [Mavis Zhou](https://www.linkedin.com/in/jiahui-mavis-zhou-a4537814a) (www.linkedin.com/in/jiahui-mavis-zhou-a4537814a)

Copyright © 2018 cognitiveclass.ai (cognitiveclass.ai?utm_source=bducopyrightlink&utm_medium=dsweb&utm_campaign=bdu). This notebook and its source code are released under the terms of the [MIT License](https://bigdatauniversity.com/mit-license/) (<https://bigdatauniversity.com/mit-license/>).