



Watson Studio democratizes machine learning and deep learning to accelerate infusion of AI in your business to drive innovation. Watson Studio provides a suite of tools and a collaborative environment for data scientists, developers and domain experts.

(http://cocl.us/pytorch_link_top)



Linear Regression 1D: Training Two Parameter Mini-Batch Gradient Descent

Table of Contents

In this lab, you will create a model the PyTorch way, this will help you as models get more complicated

- [Make Some Data](#)
- [Create the Model and Cost Function the PyTorch way](#)
- [Train the Model: Batch Gradient Descent](#)

Estimated Time Needed: **30 min**

Preparation

We'll need the following libraries:

In [1]:

```
# These are the libraries we are going to use in the lab.

import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits import mplot3d
```

The class `plot_error_surfaces` is just to help you visualize the data space and the parameter space during training and has nothing to do with PyTorch.

In [2]:

```
# class for plotting
```

```
class plot_error_surfaces(object):
```

```
    # Constructor
```

```
    def __init__(self, w_range, b_range, X, Y, n_samples = 30, go = True):
```

```
        W = np.linspace(-w_range, w_range, n_samples)
```

```
        B = np.linspace(-b_range, b_range, n_samples)
```

```
        w, b = np.meshgrid(W, B)
```

```
        Z = np.zeros((30, 30))
```

```
        count1 = 0
```

```
        self.y = Y.numpy()
```

```
        self.x = X.numpy()
```

```
        for w1, b1 in zip(w, b):
```

```
            count2 = 0
```

```
            for w2, b2 in zip(w1, b1):
```

```
                Z[count1, count2] = np.mean((self.y - w2 * self.x + b2) ** 2)
```

```
                count2 += 1
```

```
            count1 += 1
```

```
        self.Z = Z
```

```
        self.w = w
```

```
        self.b = b
```

```
        self.W = []
```

```
        self.B = []
```

```
        self.LOSS = []
```

```
        self.n = 0
```

```
        if go == True:
```

```
            plt.figure()
```

```
            plt.figure(figsize = (7.5, 5))
```

```
            plt.axes(projection = '3d').plot_surface(self.w, self.b, self.Z, rstride
```

```
e = 1, cstride = 1, cmap = 'viridis', edgecolor = 'none')
```

```
            plt.title('Loss Surface')
```

```
            plt.xlabel('w')
```

```
            plt.ylabel('b')
```

```
            plt.show()
```

```
            plt.figure()
```

```
            plt.title('Loss Surface Contour')
```

```
            plt.xlabel('w')
```

```
            plt.ylabel('b')
```

```
            plt.contour(self.w, self.b, self.Z)
```

```
            plt.show()
```

```
    # Setter
```

```
    def set_para_loss(self, model, loss):
```

```
        self.n = self.n + 1
```

```
        self.LOSS.append(loss)
```

```
        self.W.append(list(model.parameters())[0].item())
```

```
        self.B.append(list(model.parameters())[1].item())
```

```
    # Plot diagram
```

```
    def final_plot(self):
```

```
        ax = plt.axes(projection = '3d')
```

```
        ax.plot_wireframe(self.w, self.b, self.Z)
```

```
        ax.scatter(self.W, self.B, self.LOSS, c = 'r', marker = 'x', s = 200, alpha
```

```
= 1)
```

```

plt.figure()
plt.contour(self.w, self.b, self.Z)
plt.scatter(self.W, self.B, c = 'r', marker = 'x')
plt.xlabel('w')
plt.ylabel('b')
plt.show()

# Plot diagram
def plot_ps(self):
    plt.subplot(121)
    plt.ylim()
    plt.plot(self.x, self.y, 'ro', label = "training points")
    plt.plot(self.x, self.W[-1] * self.x + self.B[-1], label = "estimated line"
)

    plt.xlabel('x')
    plt.ylabel('y')
    plt.ylim((-10, 15))
    plt.title('Data Space Iteration: ' + str(self.n))
    plt.subplot(122)
    plt.contour(self.w, self.b, self.Z)
    plt.scatter(self.W, self.B, c = 'r', marker = 'x')
    plt.title('Loss Surface Contour Iteration' + str(self.n) )
    plt.xlabel('w')
    plt.ylabel('b')
    plt.show()

```

Make Some Data

Import libraries and set random seed.

In [3]:

```

# Import libraries and set random seed

import torch
from torch.utils.data import Dataset, DataLoader
torch.manual_seed(1)

```

Out[3]:

```
<torch._C.Generator at 0x7f22980ae230>
```

Generate values from -3 to 3 that create a line with a slope of 1 and a bias of -1. This is the line that you need to estimate. Add some noise to the data:

In [4]:

```
# Create Data Class

class Data(Dataset):

    # Constructor
    def __init__(self):
        self.x = torch.arange(-3, 3, 0.1).view(-1, 1)
        self.f = 1 * self.x - 1
        self.y = self.f + 0.1 * torch.randn(self.x.size())
        self.len = self.x.shape[0]

    # Getter
    def __getitem__(self, index):
        return self.x[index], self.y[index]

    # Get Length
    def __len__(self):
        return self.len
```

Create a dataset object:

In [5]:

```
# Create dataset object

dataset = Data()
```

Plot out the data and the line.


In [6]:

```
# Plot the data
```

```
plt.plot(dataset.x.numpy(), dataset.y.numpy(), 'rx', label = 'y')  
plt.plot(dataset.x.numpy(), dataset.f.numpy(), label = 'f')  
plt.xlabel('x')  
plt.ylabel('y')  
plt.legend()
```

Out[6]:

<matplotlib.legend.Legend at 0x7f22230621d0>



Create the Model and Total Loss Function (Cost)

Create a linear regression class

In [7]:

```
# Create a linear regression model class

from torch import nn, optim

class linear_regression(nn.Module):

    # Constructor
    def __init__(self, input_size, output_size):
        super(linear_regression, self).__init__()
        self.linear = nn.Linear(input_size, output_size)

    # Prediction
    def forward(self, x):
        yhat = self.linear(x)
        return yhat
```

We will use PyTorch build-in functions to create a criterion function; this calculates the total loss or cost

In [8]:

```
# Build in cost function

criterion = nn.MSELoss()
```

Create a linear regression object and optimizer object, the optimizer object will use the linear regression object.

In [9]:

```
# Create optimizer

model = linear_regression(1,1)
optimizer = optim.SGD(model.parameters(), lr = 0.01)
```


In [10]:

```
list(model.parameters())
```

Out[10]:

```
[Parameter containing:
  tensor([[0.3636]], requires_grad=True),
 Parameter containing:
  tensor([0.4957], requires_grad=True)]
```

Remember to construct an optimizer you have to give it an iterable containing the parameters i.e. provide `model.parameters()` as an input to the object constructor



Similar to the model, the optimizer has a state dictionary:

In [11]:

```
optimizer.state_dict()
```

Out[11]:

```
{'state': {},
 'param_groups': [{'lr': 0.01,
                    'momentum': 0,
                    'dampening': 0,
                    'weight_decay': 0,
                    'nesterov': False,
                    'params': [139784592575872, 139784592575008]}]}
```

Many of the keys correspond to more advanced optimizers.

Create a `Dataloader` object:

In [12]:

```
# Create Dataloader object
```

```
trainloader = DataLoader(dataset = dataset, batch_size = 1)
```

PyTorch randomly initialises your model parameters. If we use those parameters, the result will not be very insightful as convergence will be extremely fast. So we will initialise the parameters such that they will take longer to converge, i.e. look cool

In [13]:

```
# Customize the weight and bias
```

```
model.state_dict()['linear.weight'][0] = -15
model.state_dict()['linear.bias'][0] = -10
```


Create a plotting object, not part of PyTorch, just used to help visualize

In [14]:

```
# Create plot surface object  
get_surface = plot_error_surfaces(15, 13, dataset.x, dataset.y, 30, go = False)
```

Train the Model via Batch Gradient Descent

Run 10 epochs of stochastic gradient descent: **bug** data space is 1 iteration ahead of parameter space.

In [15]:


```
# Train Model

def train_model_BGD(iter):
    for epoch in range(iter):
        for x,y in trainloader:
            yhat = model(x)
            loss = criterion(yhat, y)
            get_surface.set_para_loss(model, loss.tolist())
            optimizer.zero_grad()
            loss.backward()


            optimizer.step()
            get_surface.plot_ps()

train_model_BGD(10)
```


Data Space Iteration: 60 Loss Surface Contour Iteration60




Data Space Iteration: 120 Loss Surface Contour Iteration120




Data Space Iteration: 180 Loss Surface Contour Iteration180



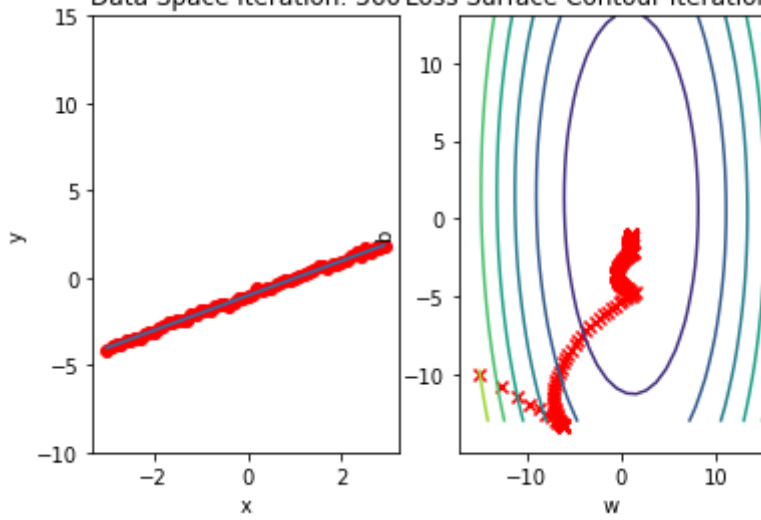
Data Space Iteration: 240 Loss Surface Contour Iteration240




Data Space Iteration: 300 Loss Surface Contour Iteration300




Data Space Iteration: 360 Loss Surface Contour Iteration360




Data Space Iteration: 420 Loss Surface Contour Iteration420

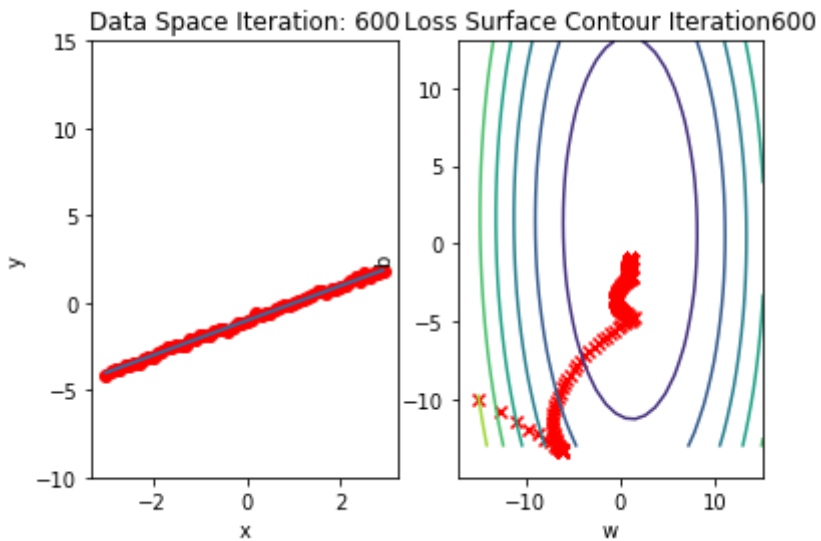


Data Space Iteration: 480 Loss Surface Contour Iteration480



Data Space Iteration: 540 Loss Surface Contour Iteration540






In [16]:

```
model.state_dict()
```

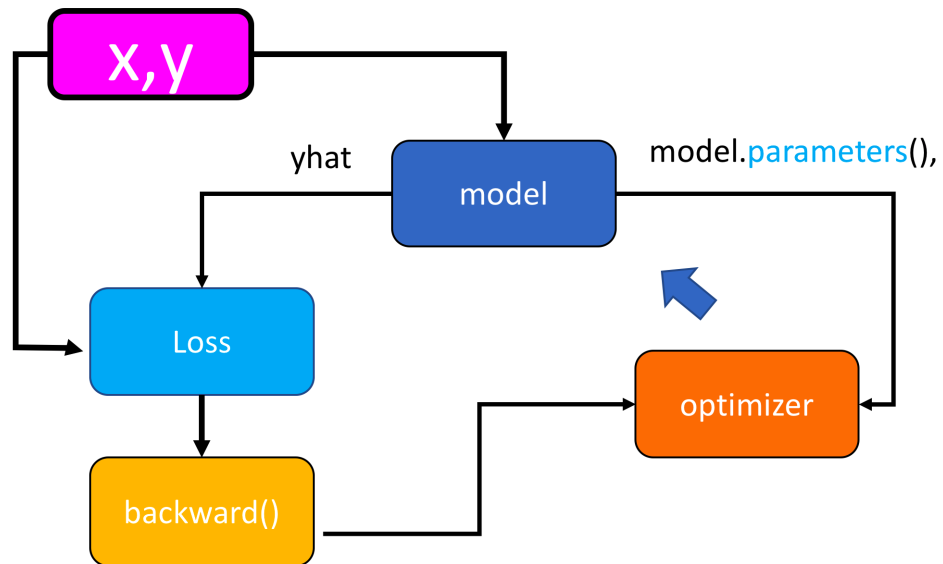
Out[16]:

```
OrderedDict([('linear.weight', tensor([[0.9932]])),
            ('linear.bias', tensor([-1.0174]))])
```

Let's use the following diagram to help clarify the process. The model takes x to produce an estimate \hat{y} , it will then be compared to the actual y with the loss function.



When we call `backward()` on the loss function, it will handle the differentiation. Calling the method step on the optimizer object it will update the parameters as they were inputs when we constructed the optimizer object. The connection is shown in the following figure :



Practice

Try to train the model via BGD with `lr = 0.1`. Use `optimizer` and the following given variables.

In [17]:

```
# Practice: Train the model via BGD using optimizer


model = linear_regression(1,1)
model.state_dict()['linear.weight'][0] = -15
model.state_dict()['linear.bias'][0] = -10
get_surface = plot_error_surfaces(15, 13, dataset.x, dataset.y, 30, go = False)

optimizer = optim.SGD(model.parameters(), lr = 0.1)
trainloader = DataLoader(dataset = dataset, batch_size = 1)


def my_train_model(iter):
    for epoch in range(iter):
        for x,y in trainloader:
            yhat = model(x)
            loss = criterion(yhat, y)
            get_surface.set_para_loss(model, loss.tolist())
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
        get_surface.plot_ps()

my_train_model(10)
```



Data Space Iteration: 60 Loss Surface Contour Iteration60




Data Space Iteration: 120 Loss Surface Contour Iteration120




Data Space Iteration: 180 Loss Surface Contour Iteration180




Data Space Iteration: 240 Loss Surface Contour Iteration240



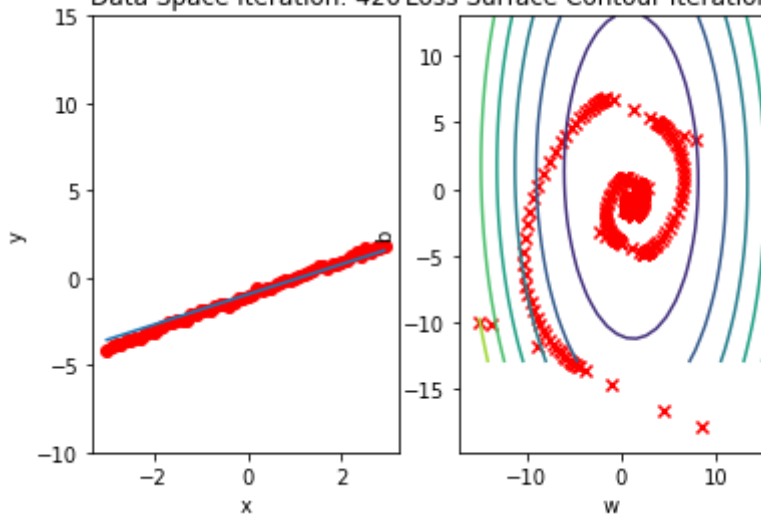
Data Space Iteration: 300 Loss Surface Contour Iteration300




Data Space Iteration: 360 Loss Surface Contour Iteration360




Data Space Iteration: 420 Loss Surface Contour Iteration420




Data Space Iteration: 480 Loss Surface Contour Iteration480



Data Space Iteration: 540 Loss Surface Contour Iteration540







Double-click **here** for the solution.

Get IBM Watson Studio free of charge!

Build and train AI & machine learning models, prepare and analyze data – all in a flexible, hybrid cloud environment. Get IBM Watson Studio Lite Plan free of charge.



Learn
Get started or get better with built-in learning.



Create
Use the best of open source tooling with IBM innovation.



Collaborate
Work smarter using community, work faster with your team.

[Sign Up For a Free Trial](#)

(http://cocl.us/pytorch_link_bottom)

About the Authors:

[Joseph Santarcangelo](https://www.linkedin.com/in/joseph-s-50398b136/) (<https://www.linkedin.com/in/joseph-s-50398b136/>) has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Other contributors: [Michelle Carey](https://www.linkedin.com/in/michelleccarey/) (<https://www.linkedin.com/in/michelleccarey/>), [Mavis Zhou](https://www.linkedin.com/in/jiahui-mavis-zhou-a4537814a) (www.linkedin.com/in/jiahui-mavis-zhou-a4537814a)

Copyright © 2018 [cognitiveclass.ai](https://cognitiveclass.ai?utm_source=bducopyrightlink&utm_medium=dswb&utm_campaign=bdu) (cognitiveclass.ai?utm_source=bducopyrightlink&utm_medium=dswb&utm_campaign=bdu). This notebook and its source code are released under the terms of the [MIT License](https://bigdatauniversity.com/mit-license/) (<https://bigdatauniversity.com/mit-license/>).